

BIG DATA SEMINAR

TOPIC: ML-1



Keras

Group CQ2021-01

GROUP CQ2021-01



Huỳnh Cao Khôi
21120275



Phạm Lê Tú Nhi
21120308

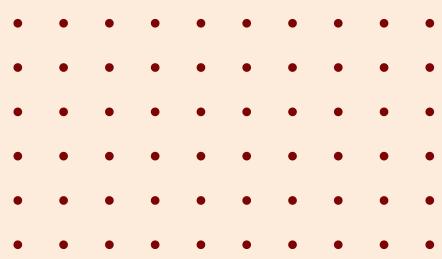
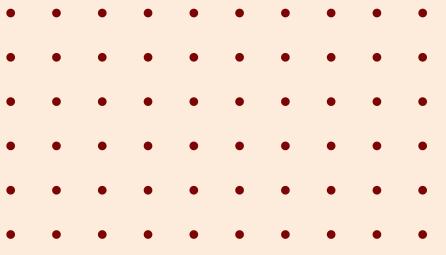


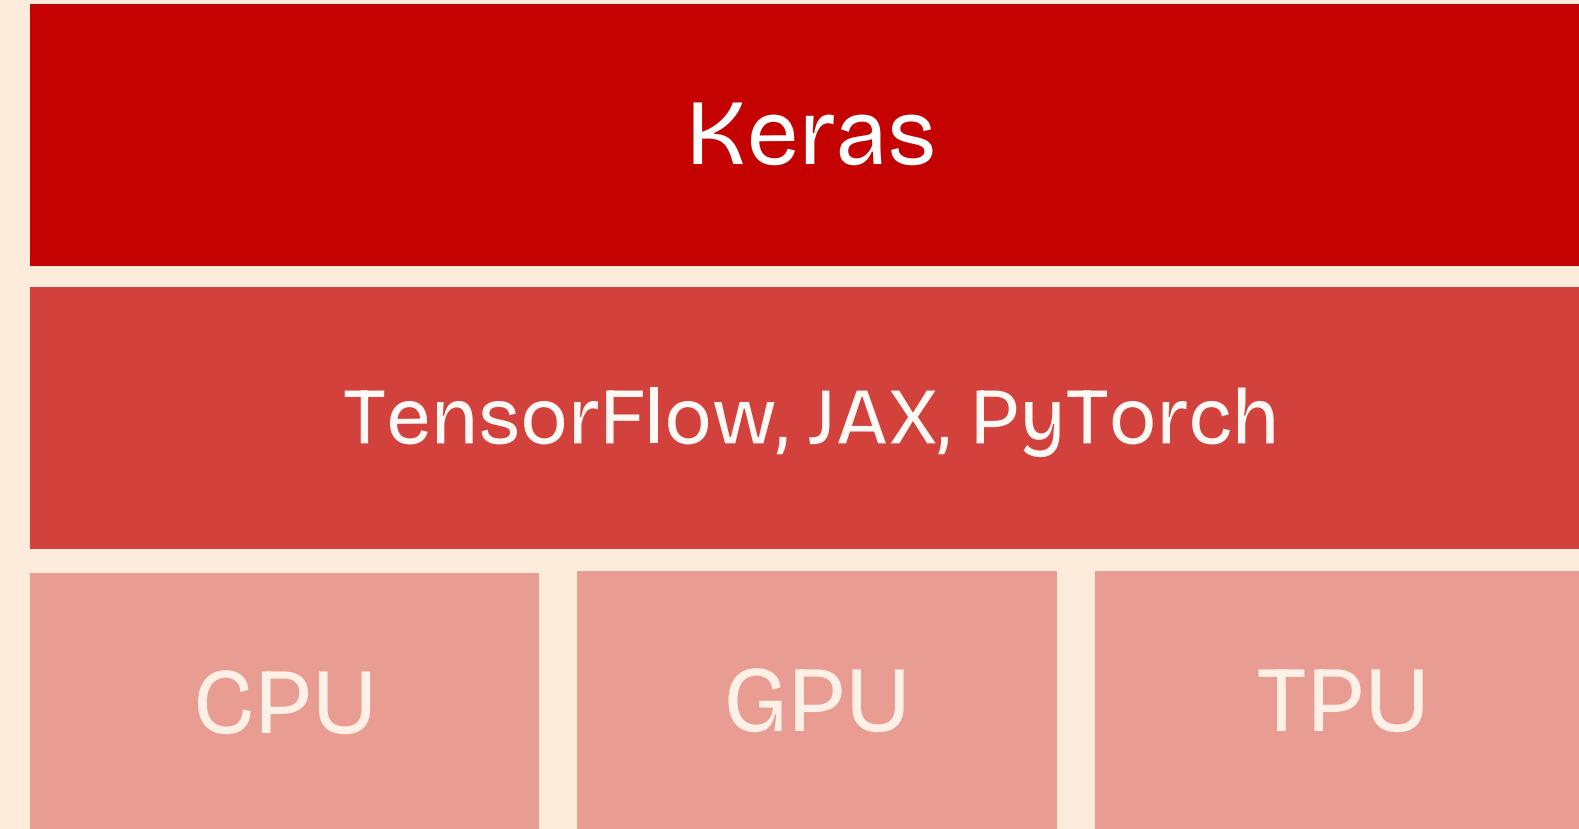
Chu Hải Linh
21120496



Lê Thị Minh Phương
21120533

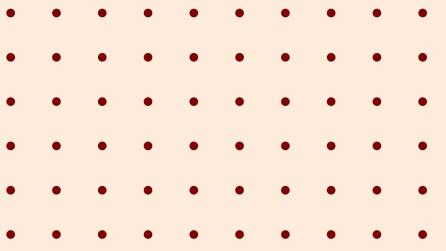
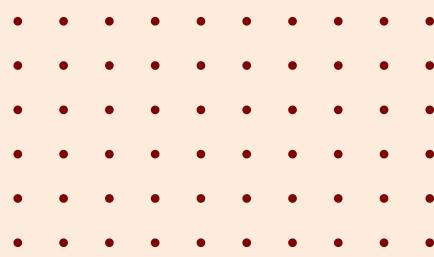
What is Keras?

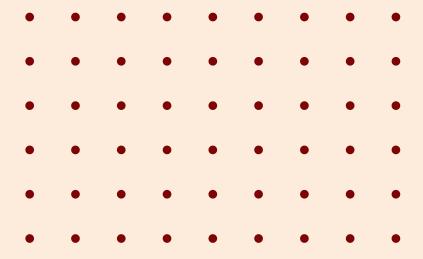




Keras is a **high-level** and **cross-backend** deep learning framework.

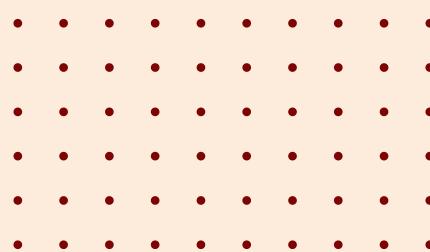
But Keras definition is **NOT**
the objective of this presentation.

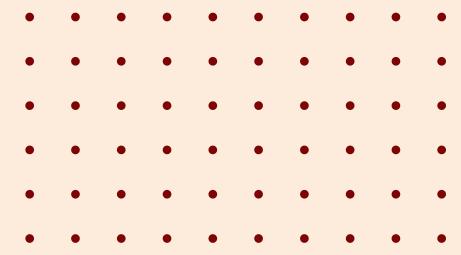




Our final objective is to show you

HOW TO USE KERAS



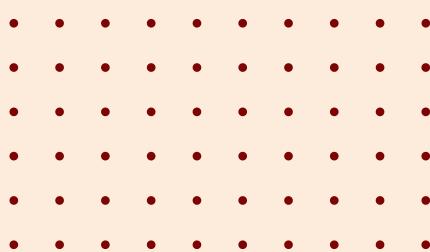


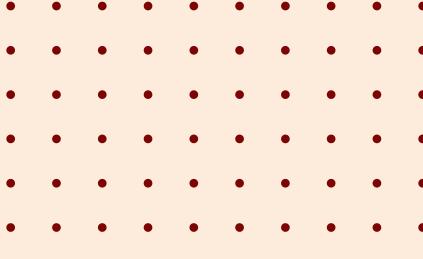
What is Keras used for?

Keras is high-level, cross-backend

deep learning

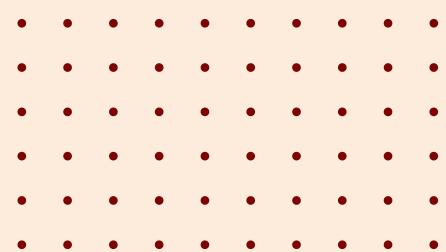
API running on top of lower APIs.





But we might have little idea what **deep learning** is...

... which makes it hard to understand what Keras is doing.



So we will approach Keras with a
“Hello World” problem in Deep Learning



MNIST

handwritten digit classification

CONTENTS

01

Problem

02

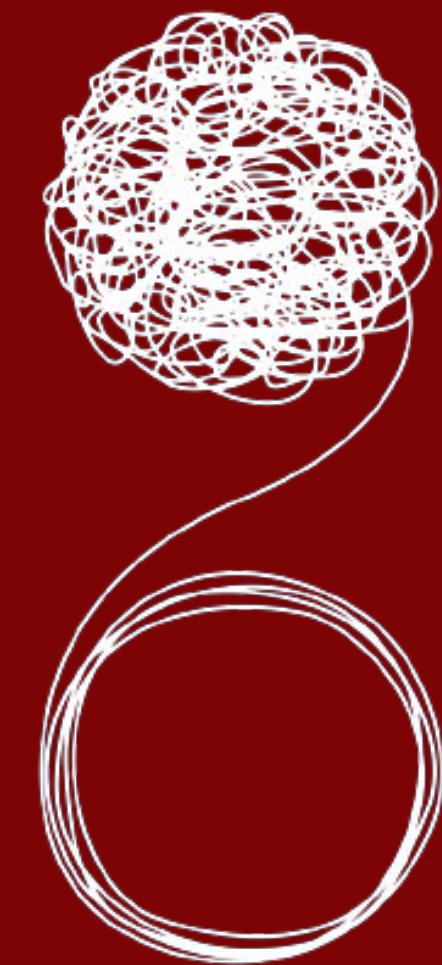
Keras

03

Demo

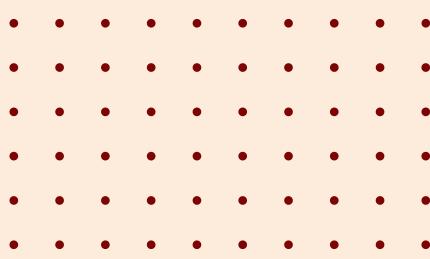
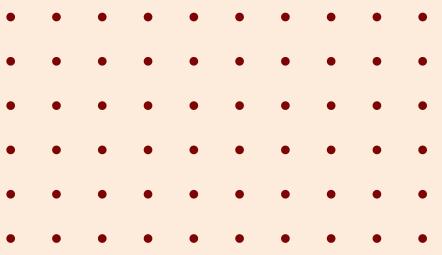
04

Conclusion



01

Problem



Handwritten digit

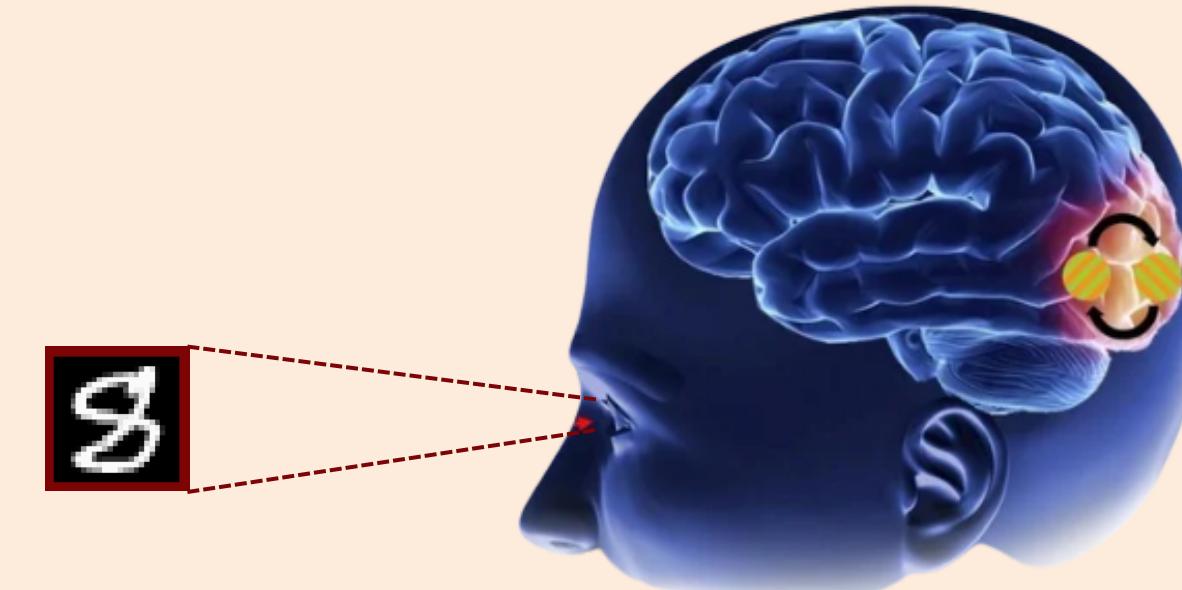


MNIST dataset

Can you classify the digit depicted in the image?

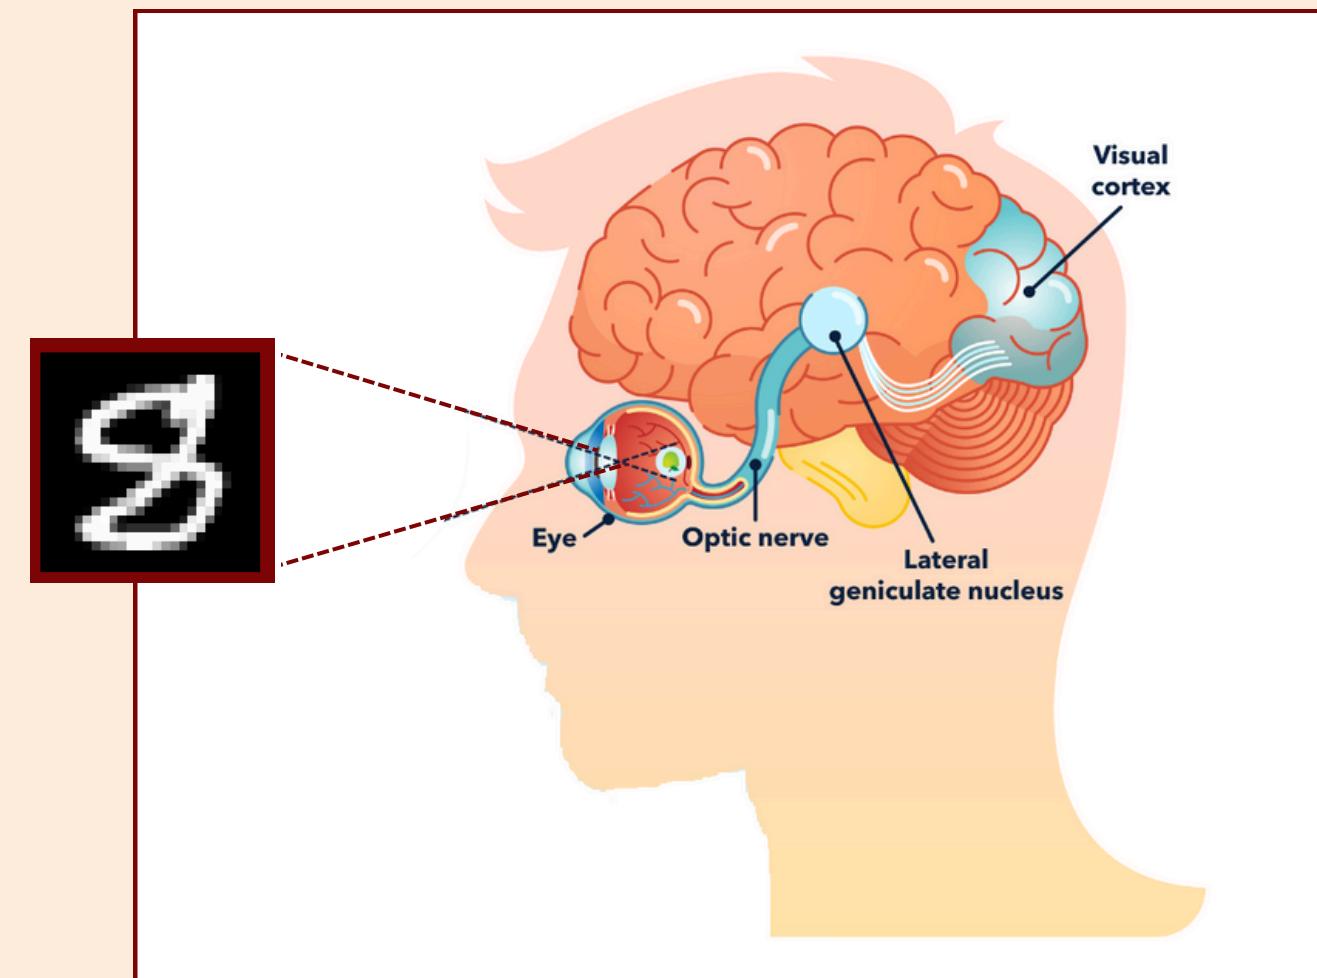


We may classify the image in the blink of an eye...

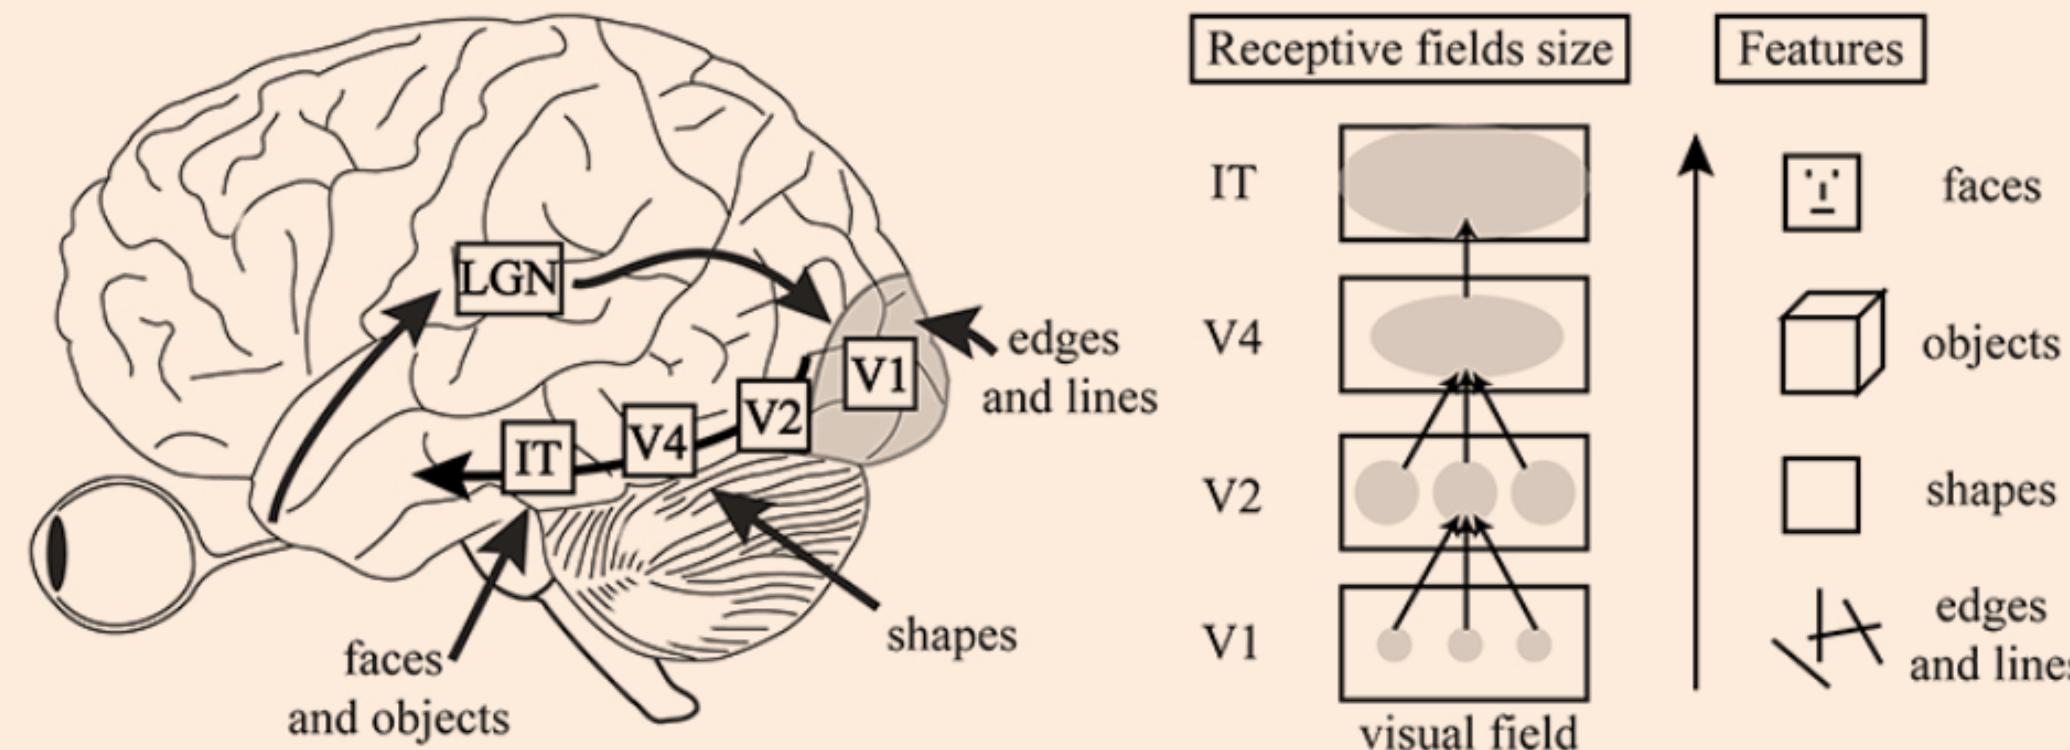


... but the way the brain does it behind the scenes
is very complicated and fascinating!

Visual information is processed in the **visual cortex**.



Features are extracted **hierarchically**, in local **receptive fields**.



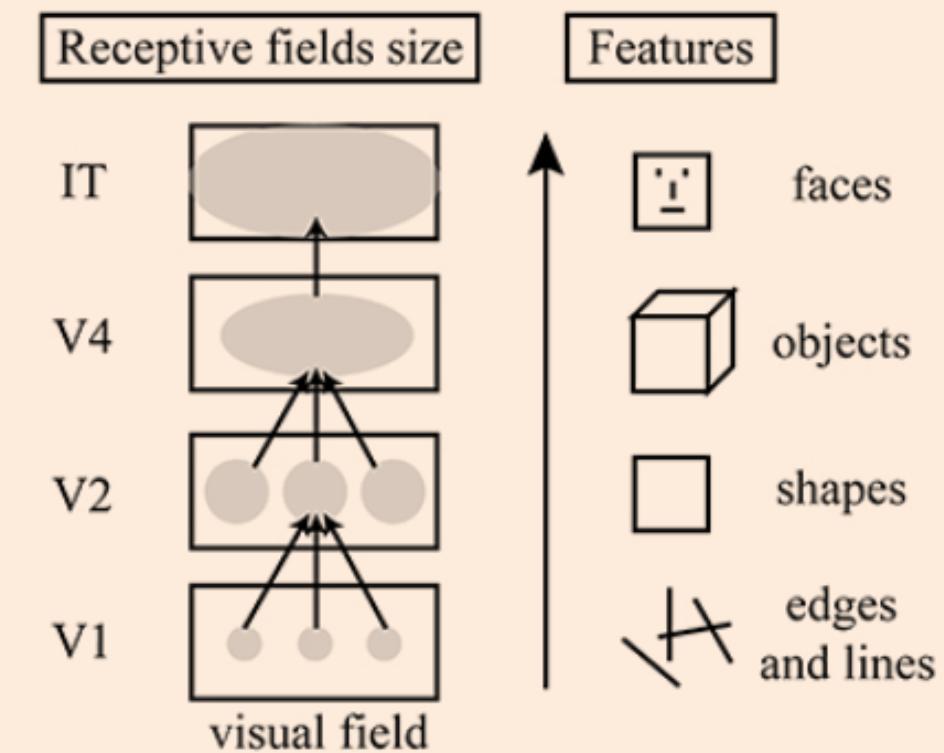
It has inspired the development of CNN

Convolutional Neural Network,

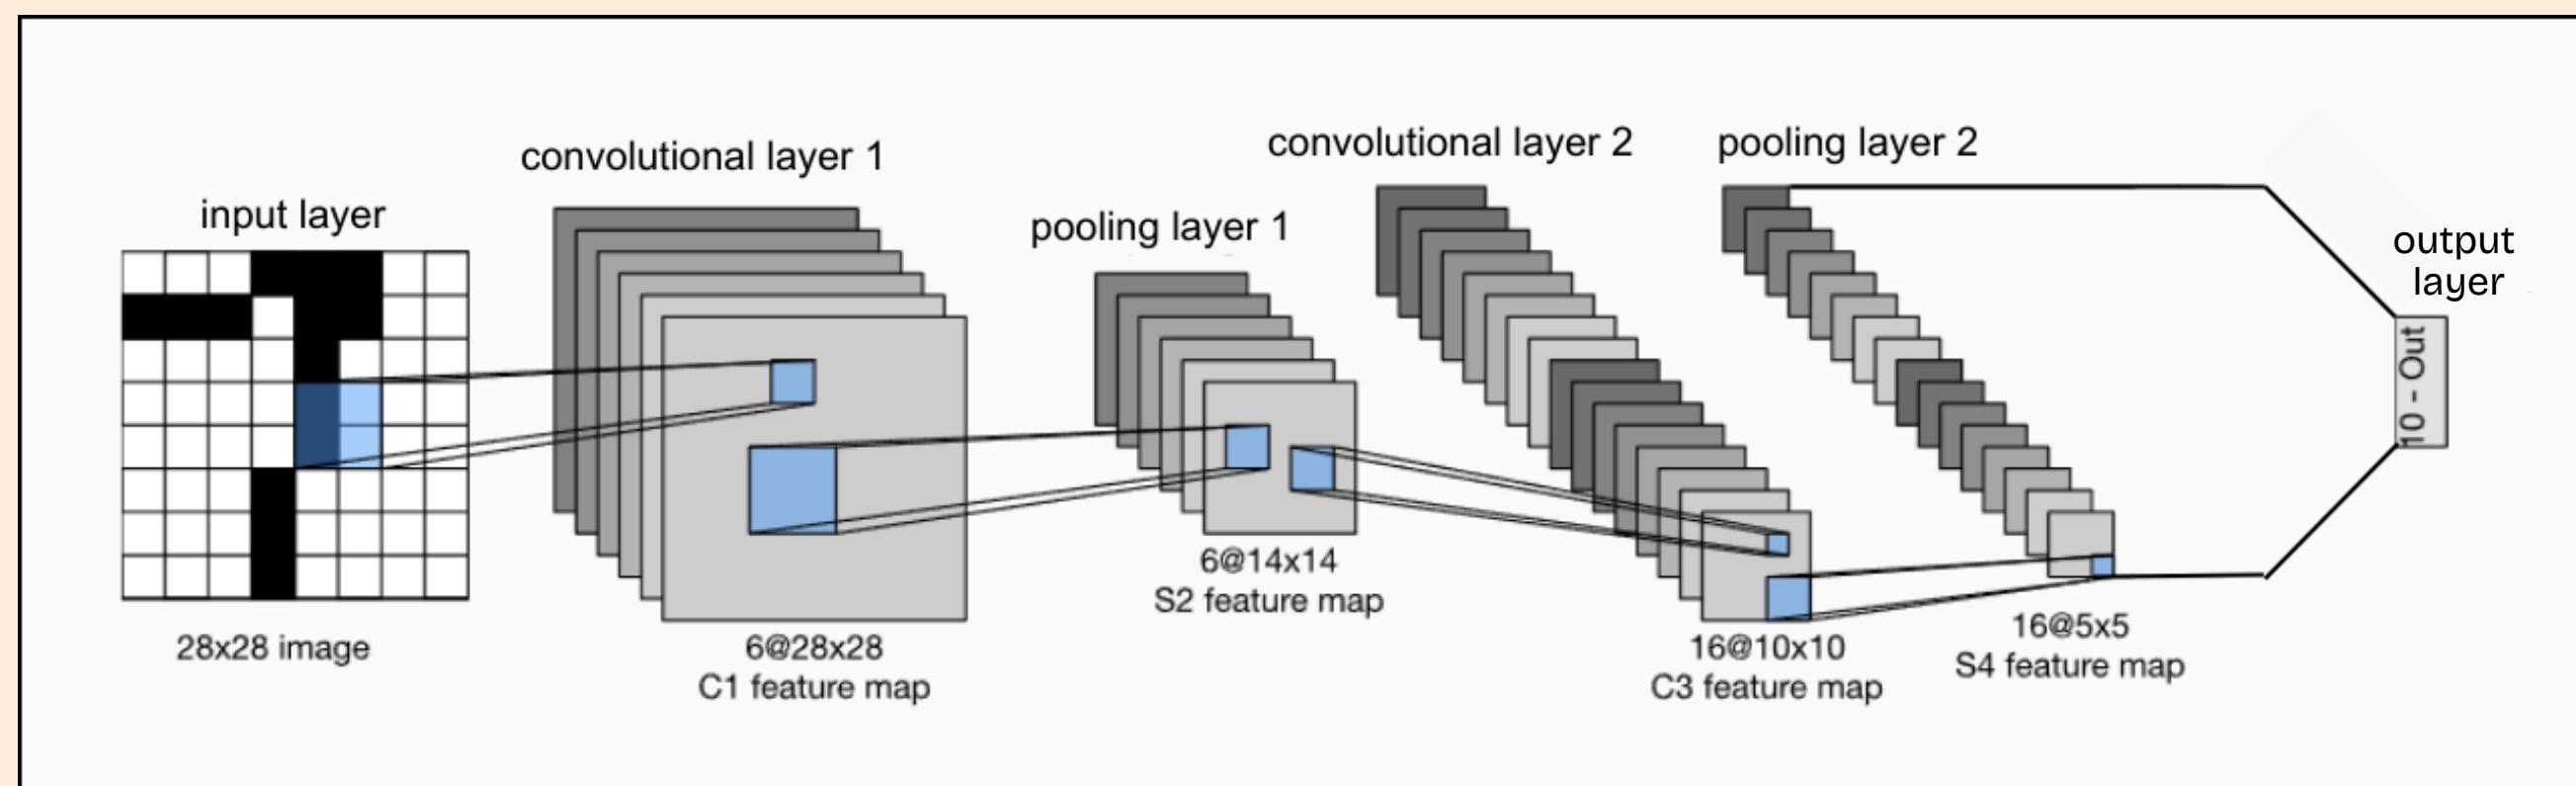
a deep learning model that performs
very well in **visual processing** tasks.

CNNs mimic how the visual cortex processes and recognizes visual patterns.

Local
Receptive
Fields

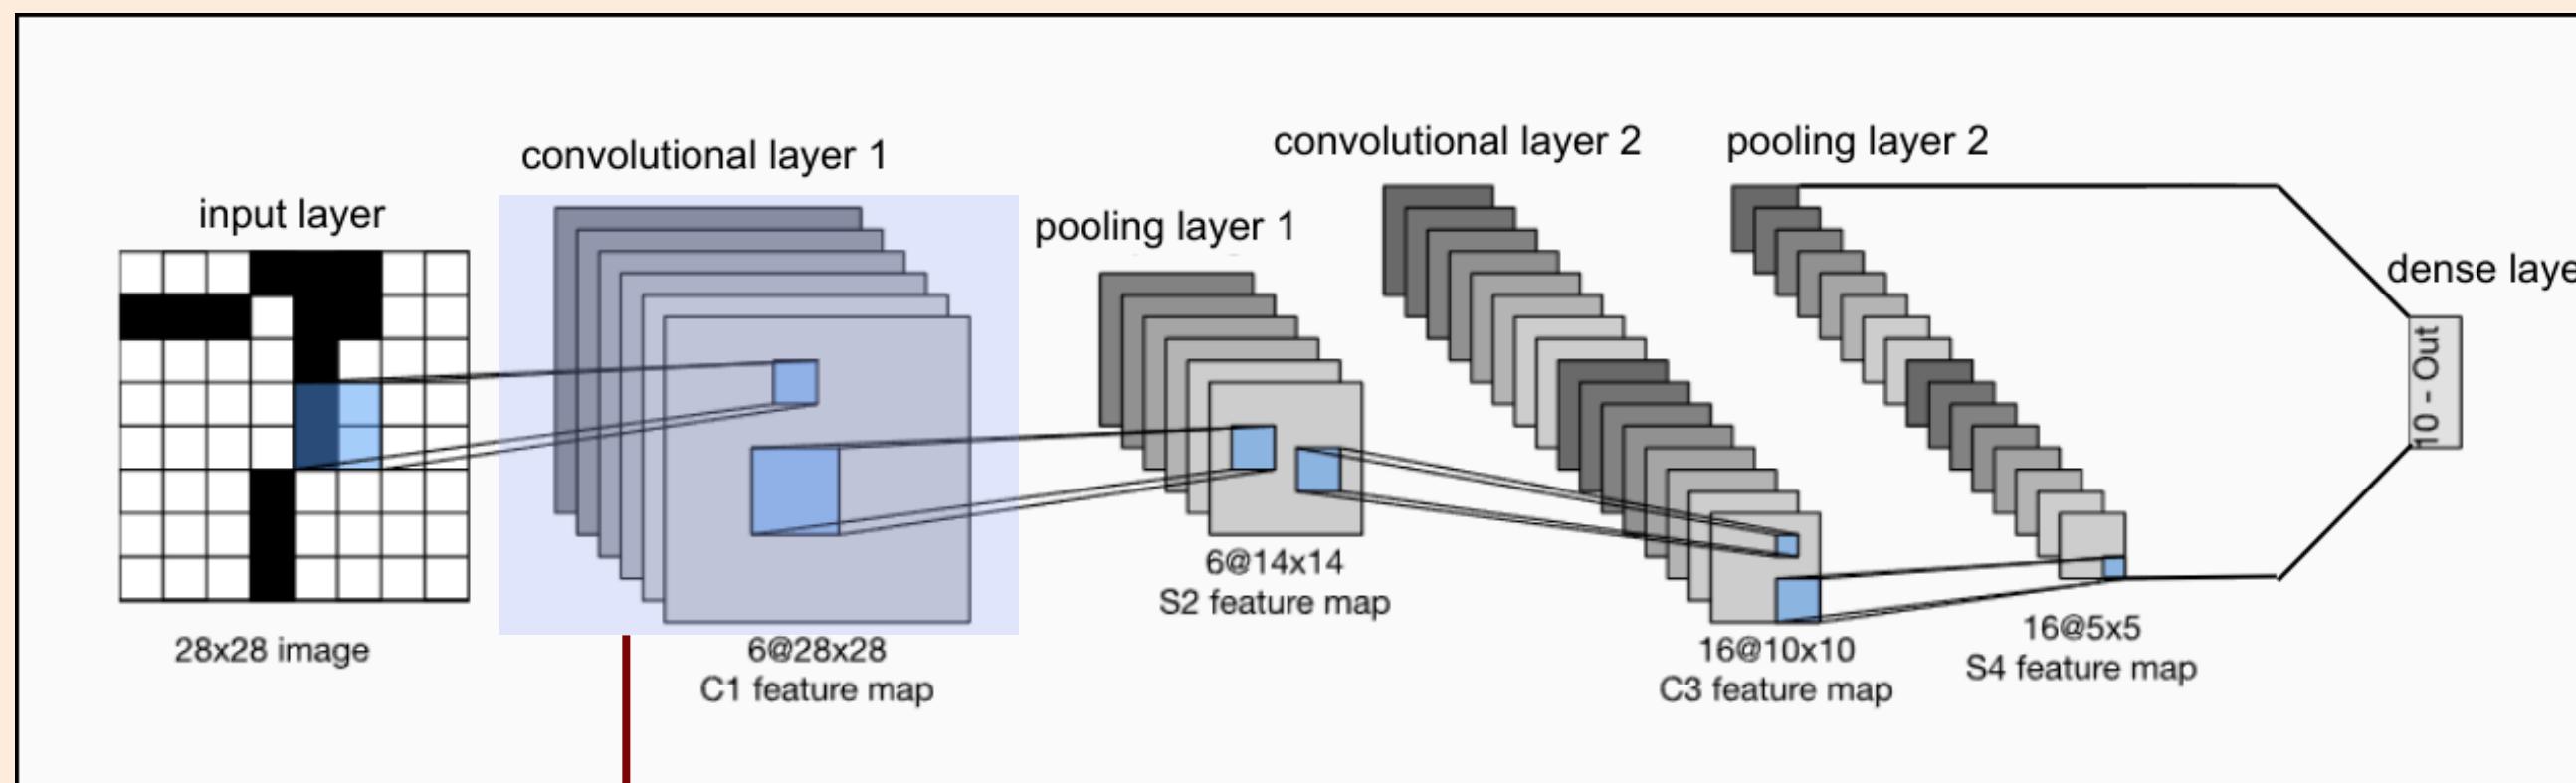


Hierarchical
Processing



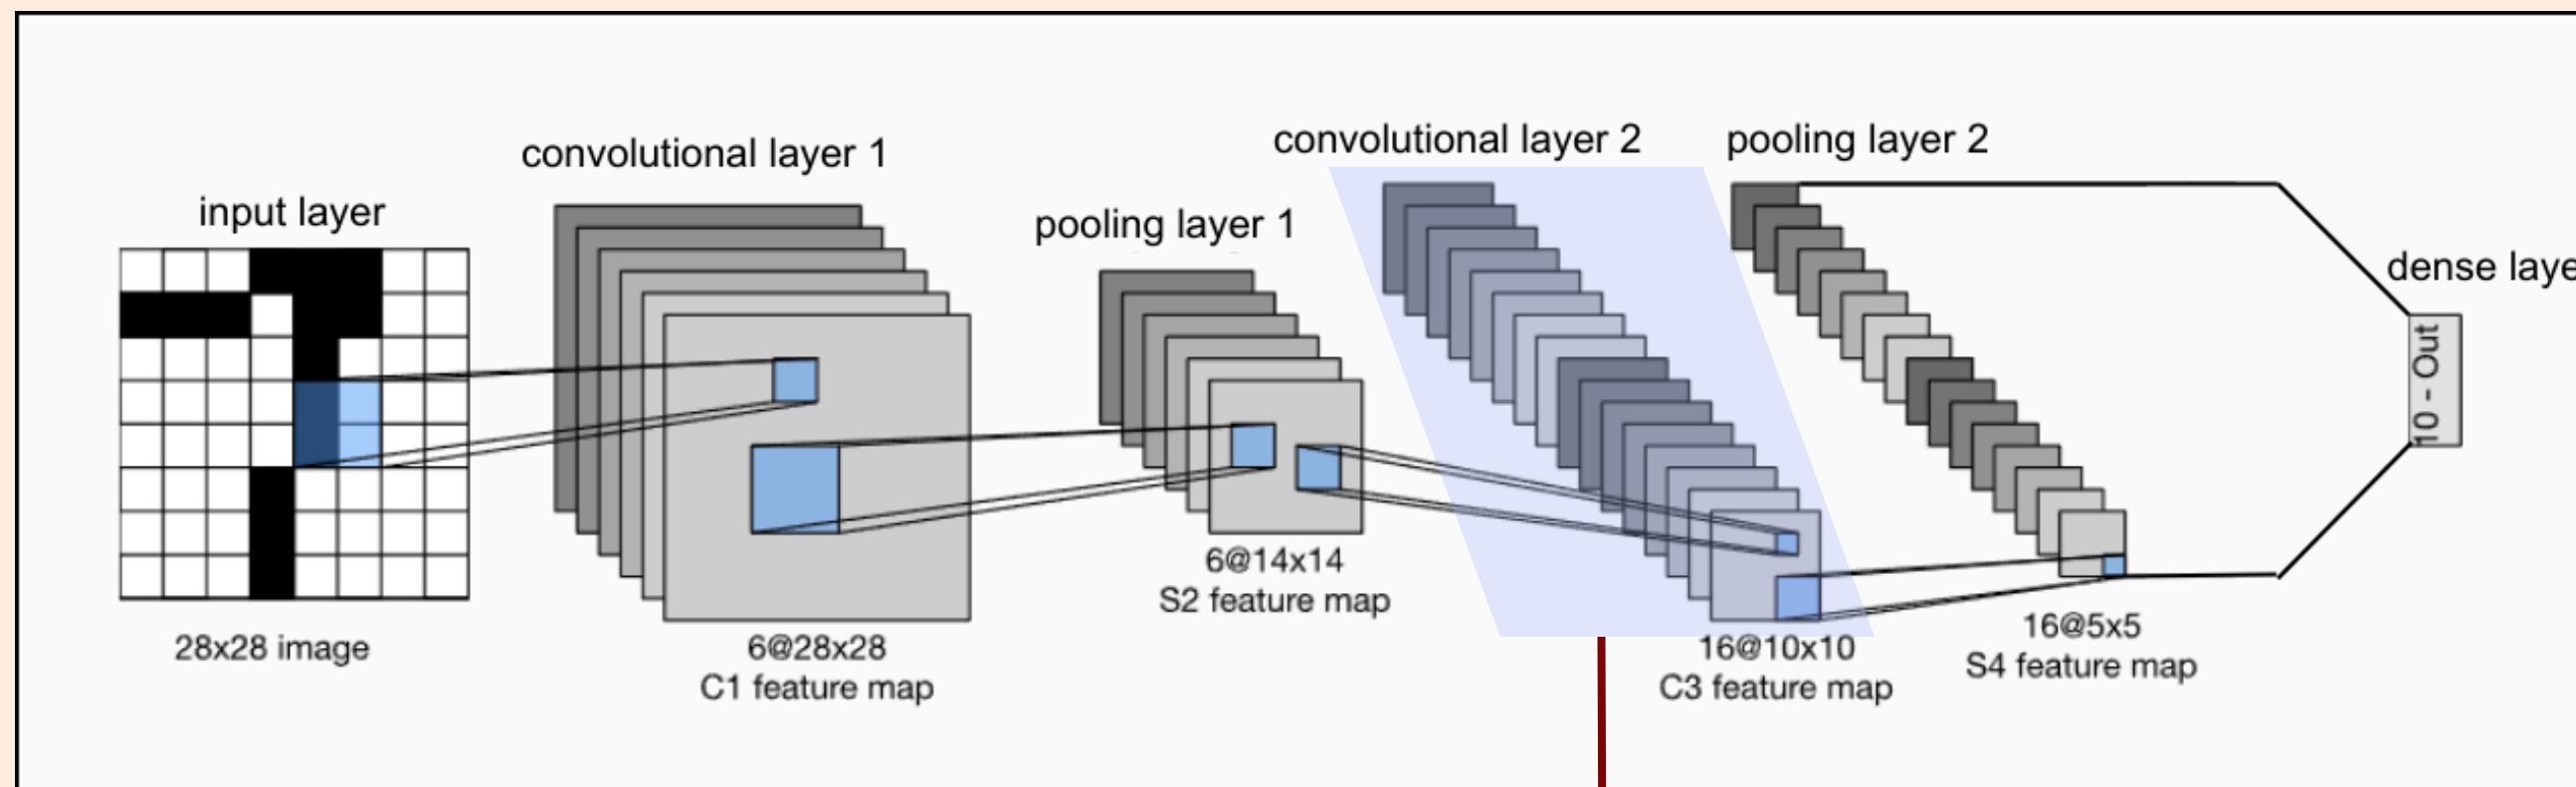
This is a **Convolutional Neural Network**.

Low-level feature extraction



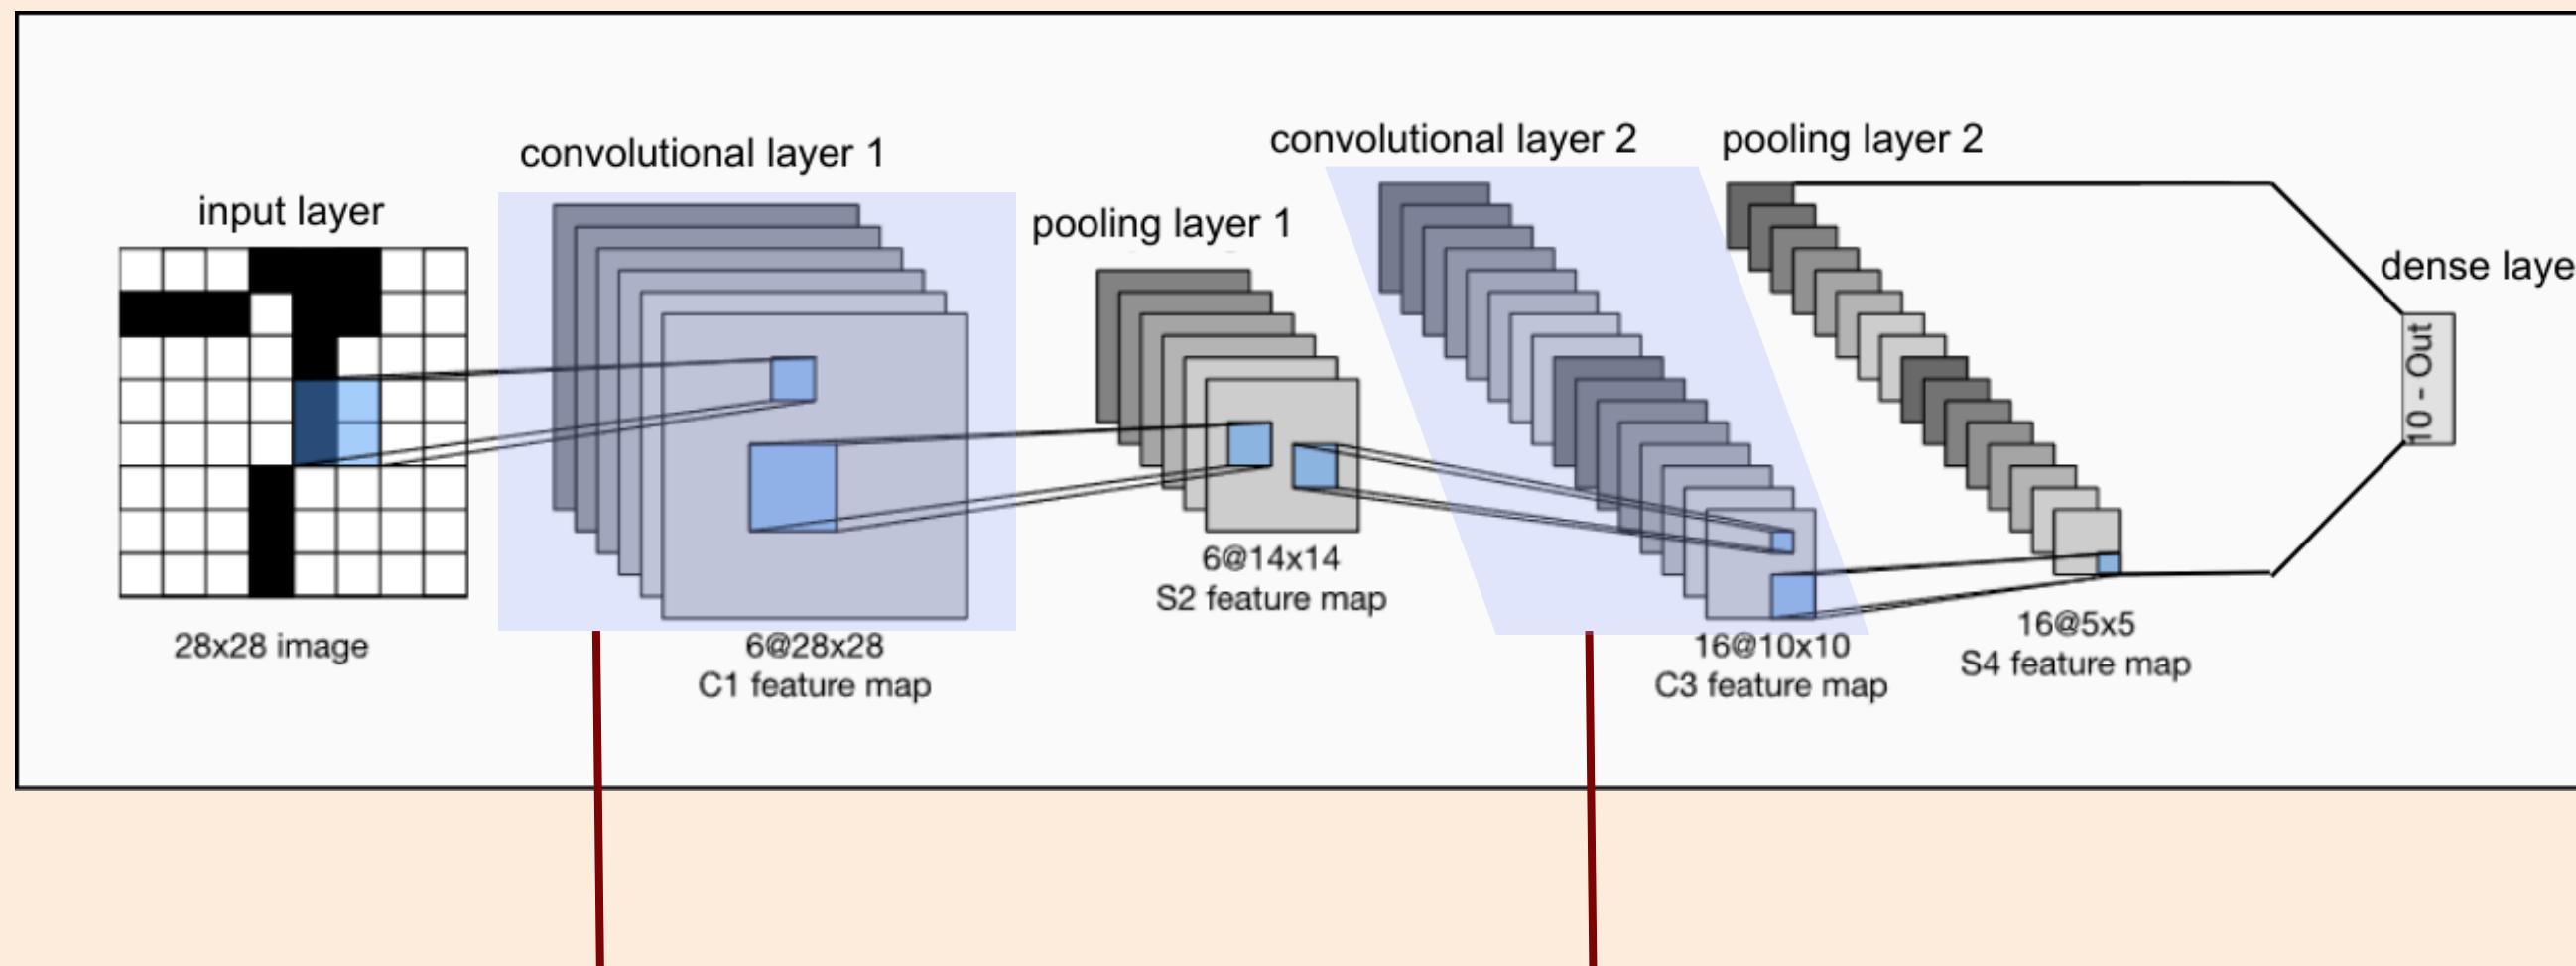
In this layer, **low-level** features (edges like vertical, horizontal and diagonal edges) are extracted.

High-level feature extraction



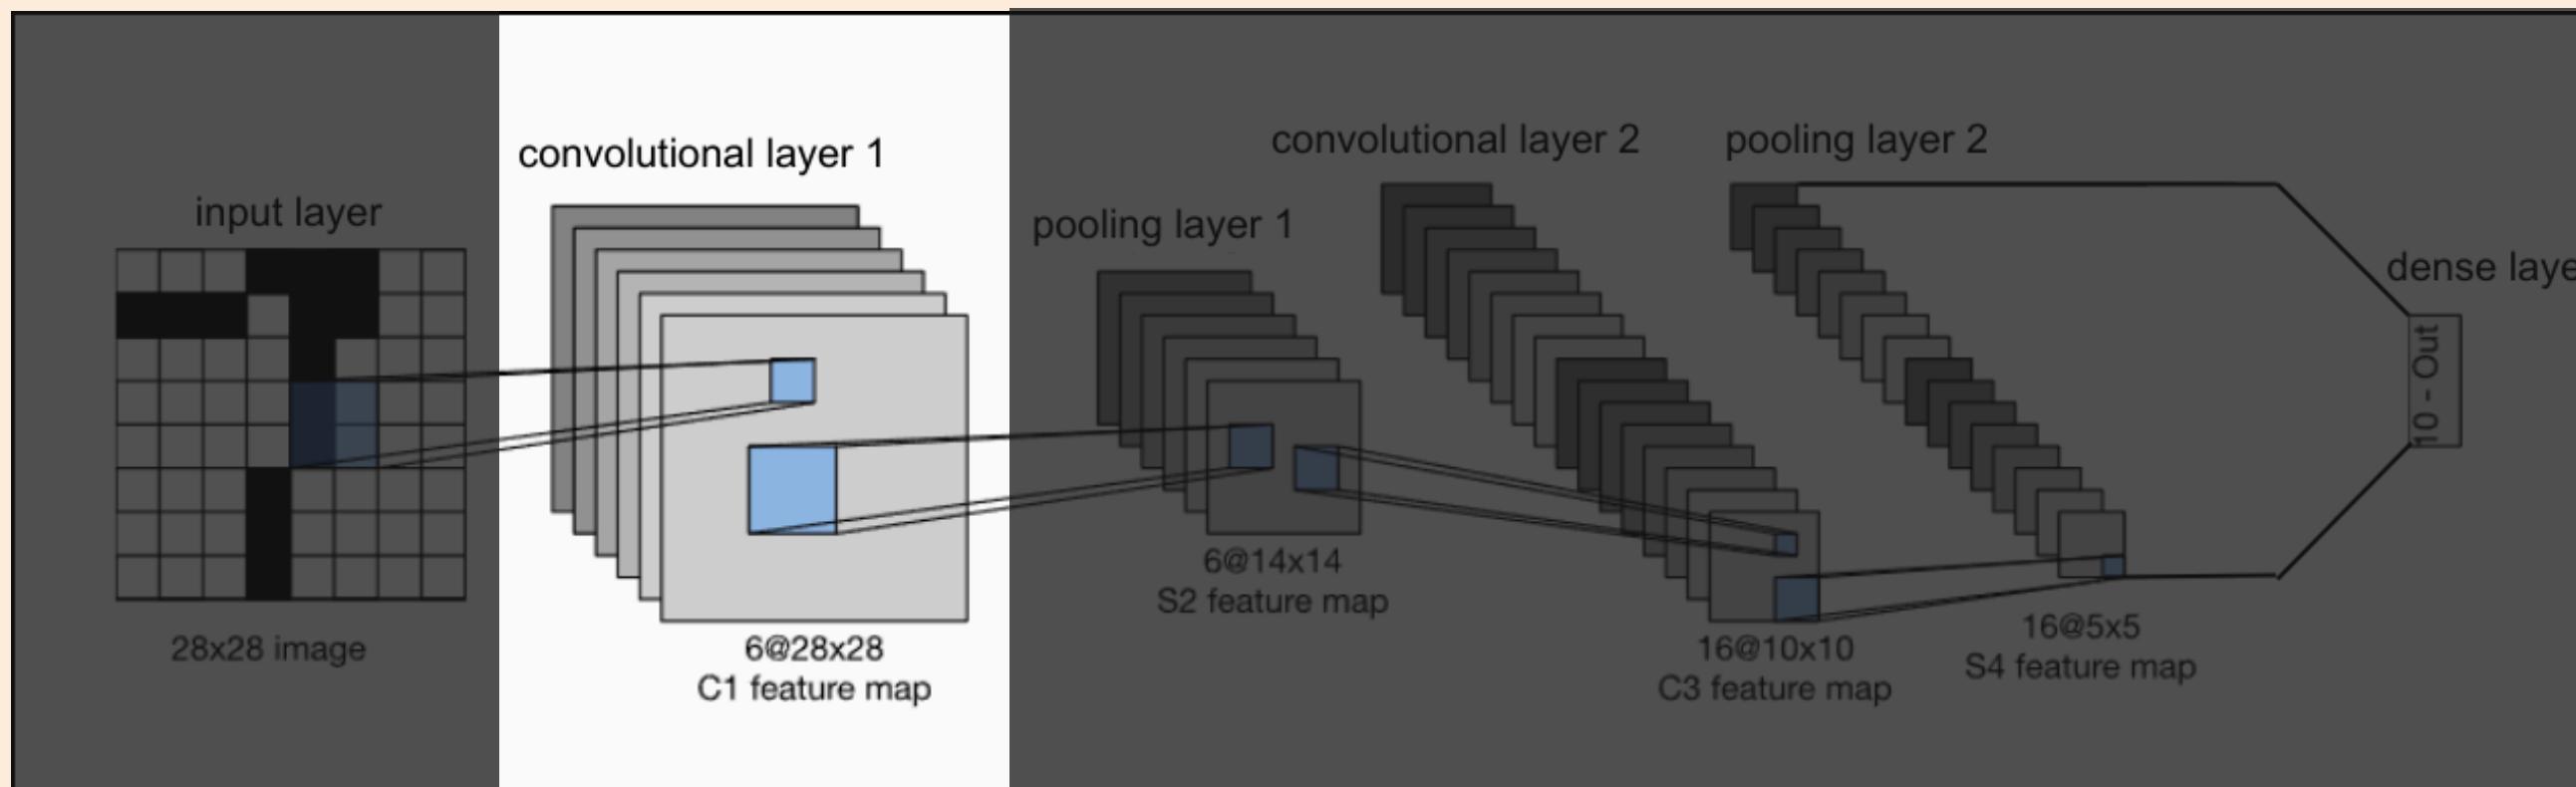
And in this layer, **higher-level** features
(shapes like circle, curve and line) are extracted.

Hierarchical feature extraction

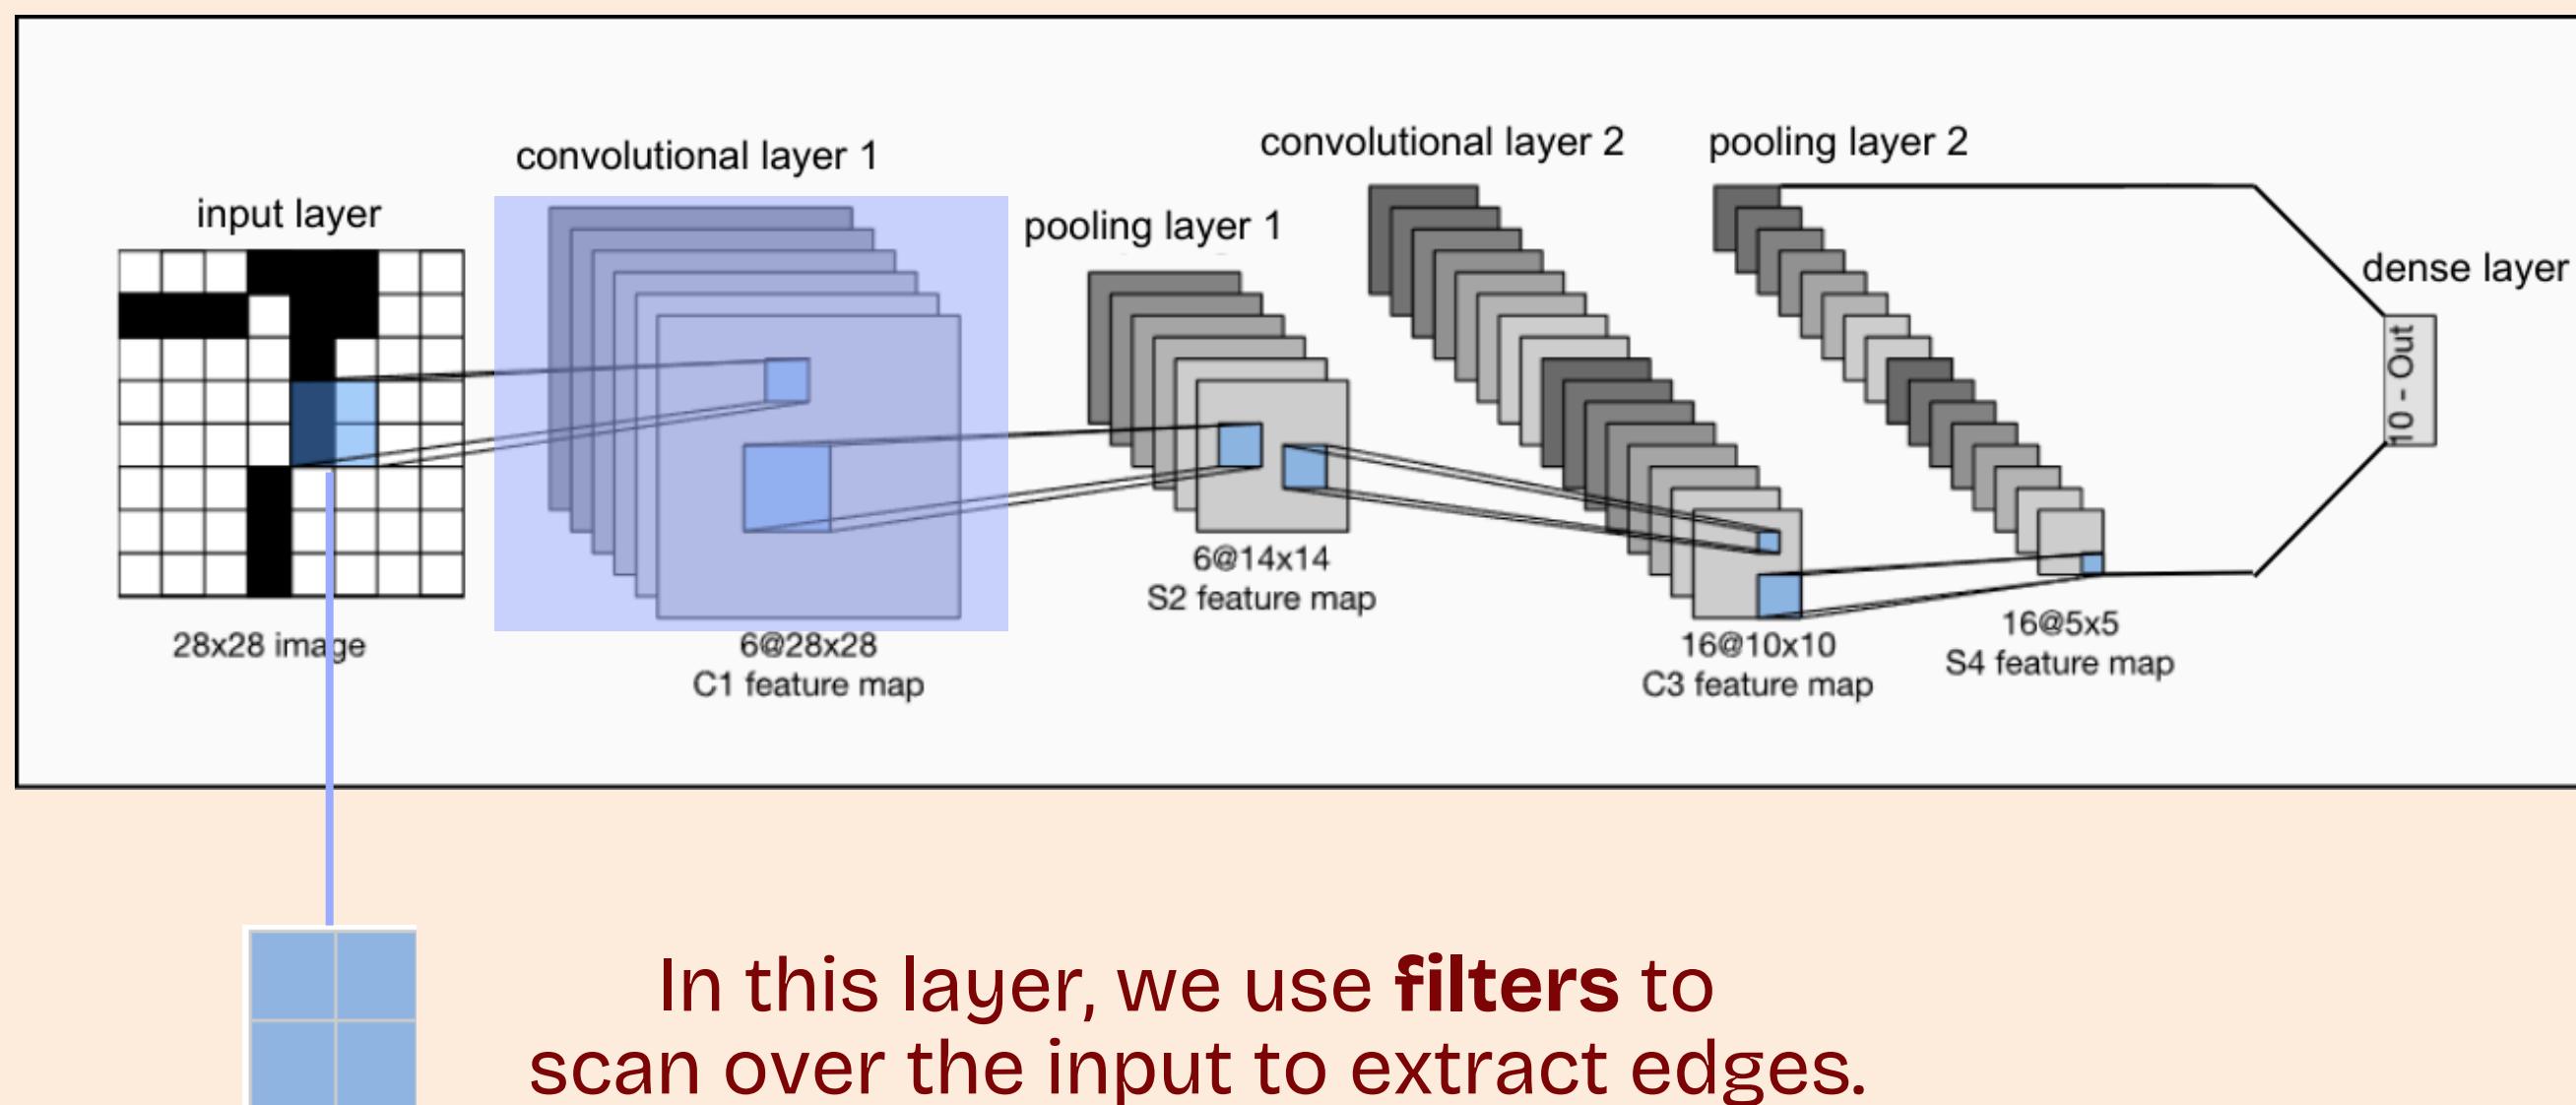


Features of the image are extracted **hierarchically**, from low to high level.

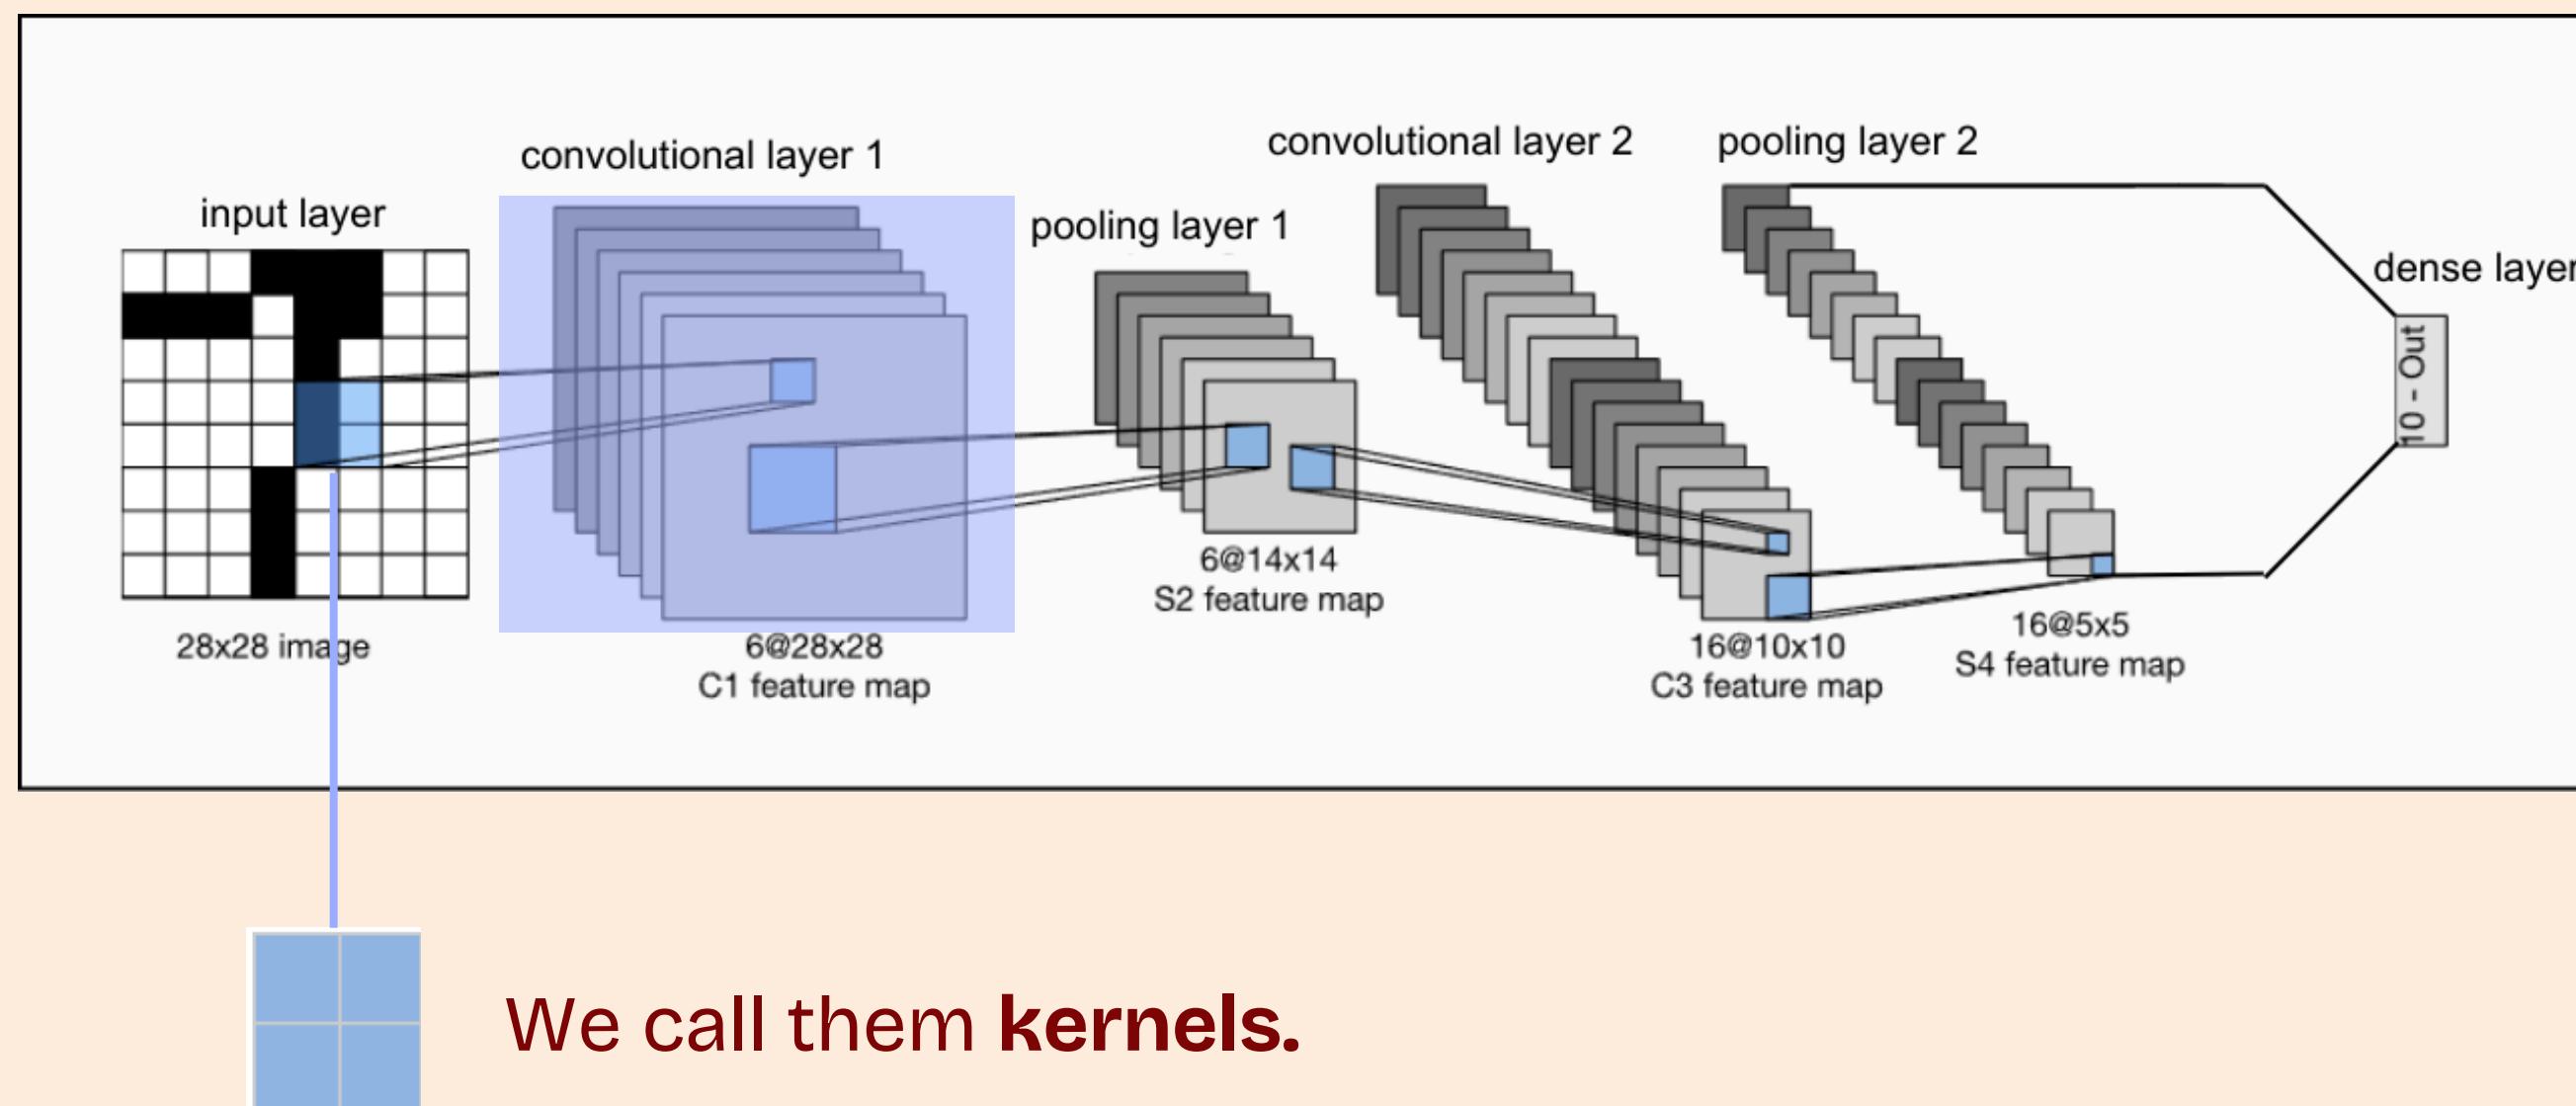
Let's zoom in this layer to understand how the model extract features.



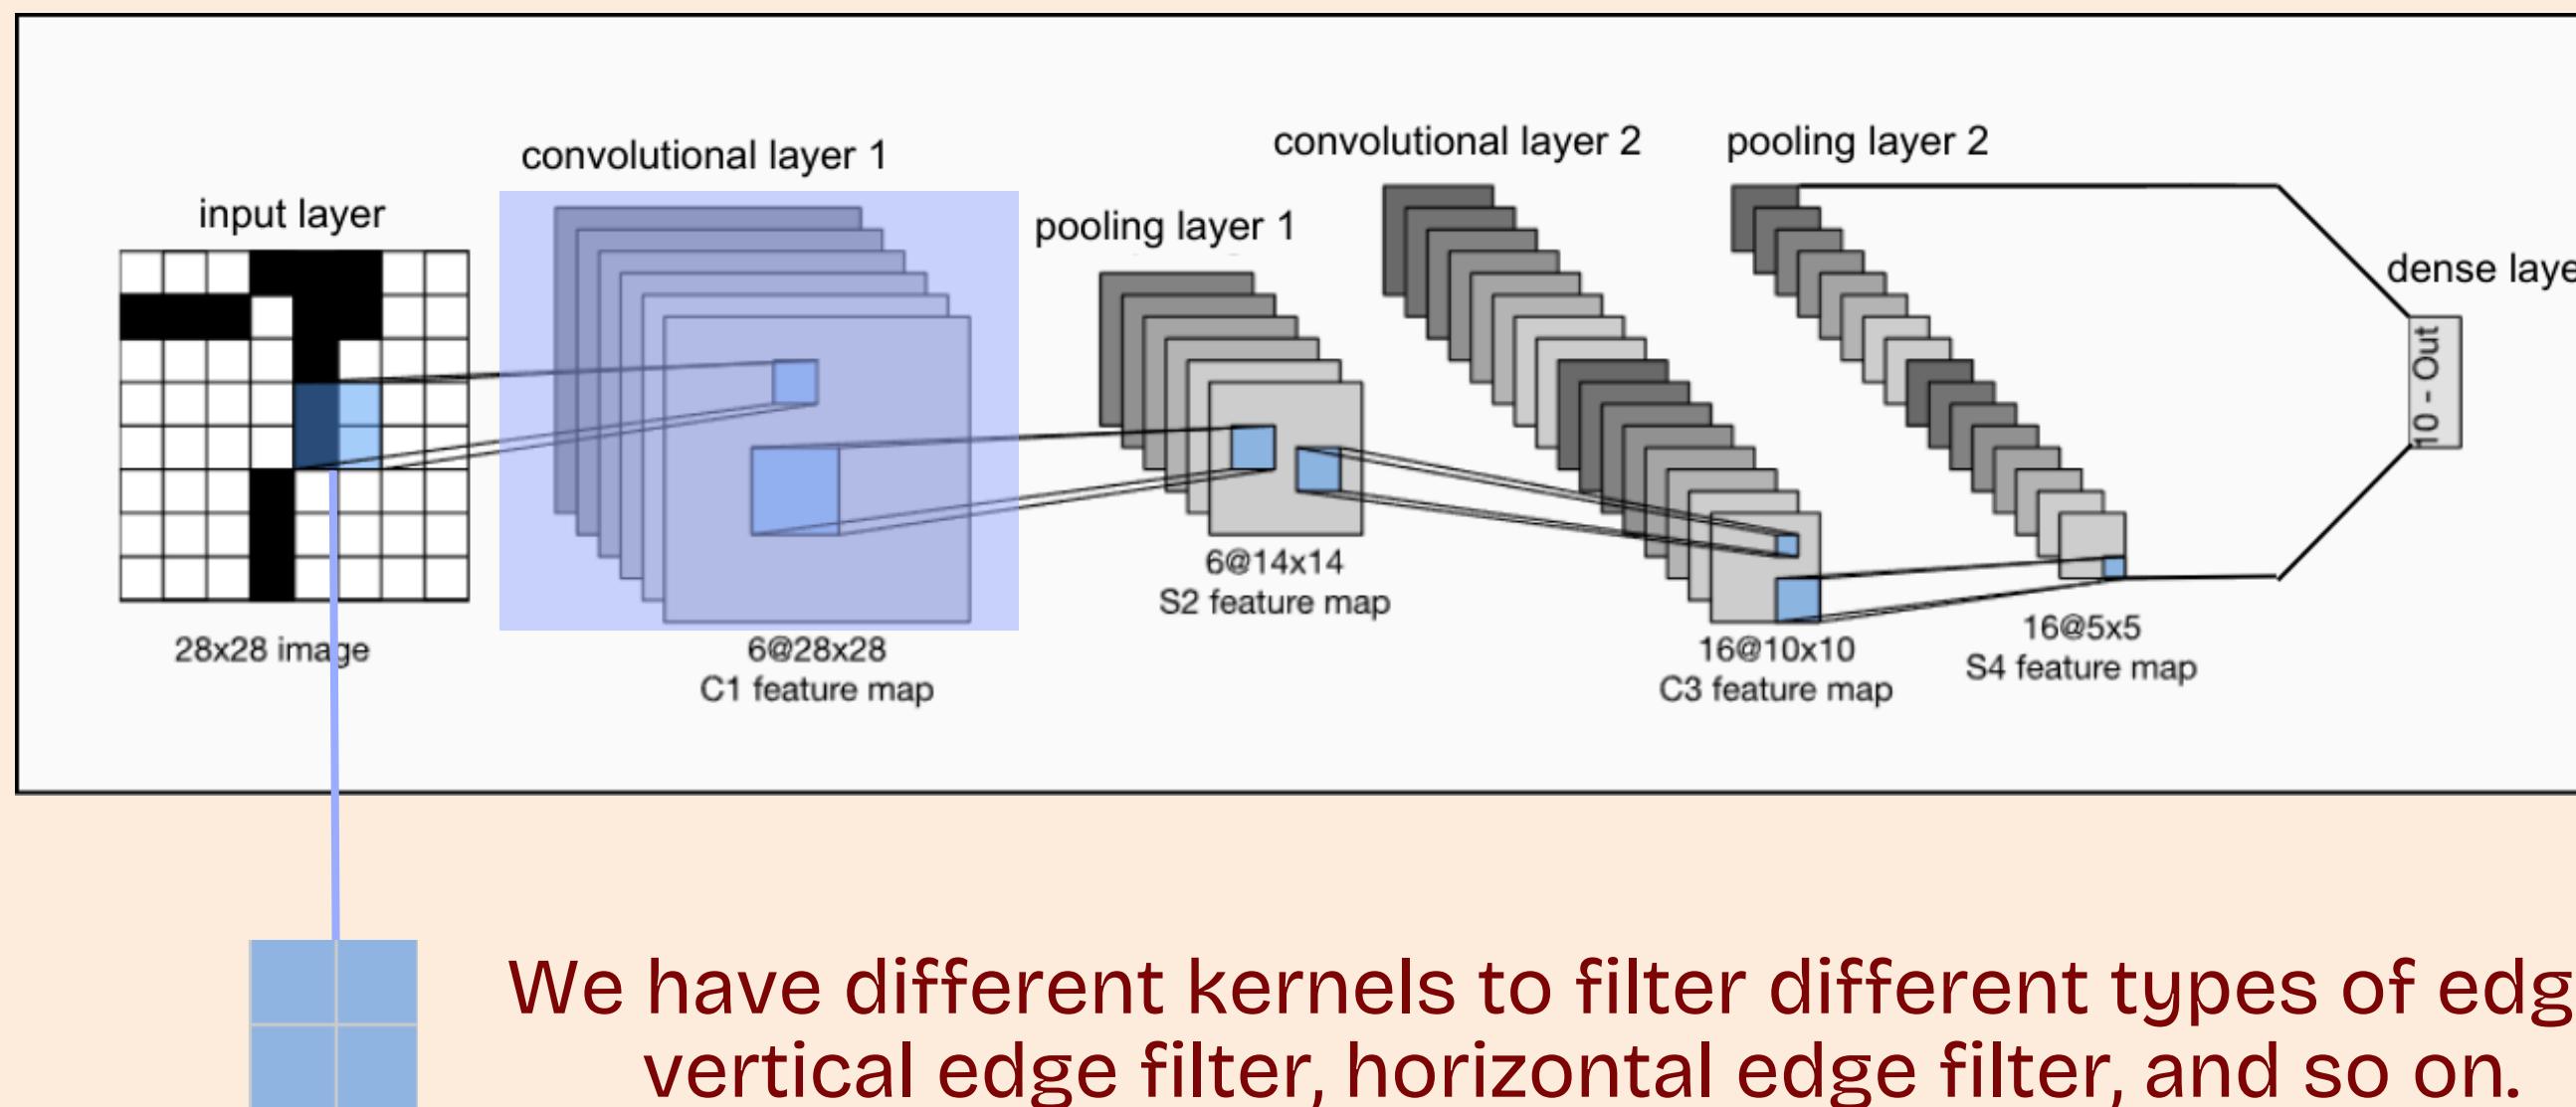
Filter



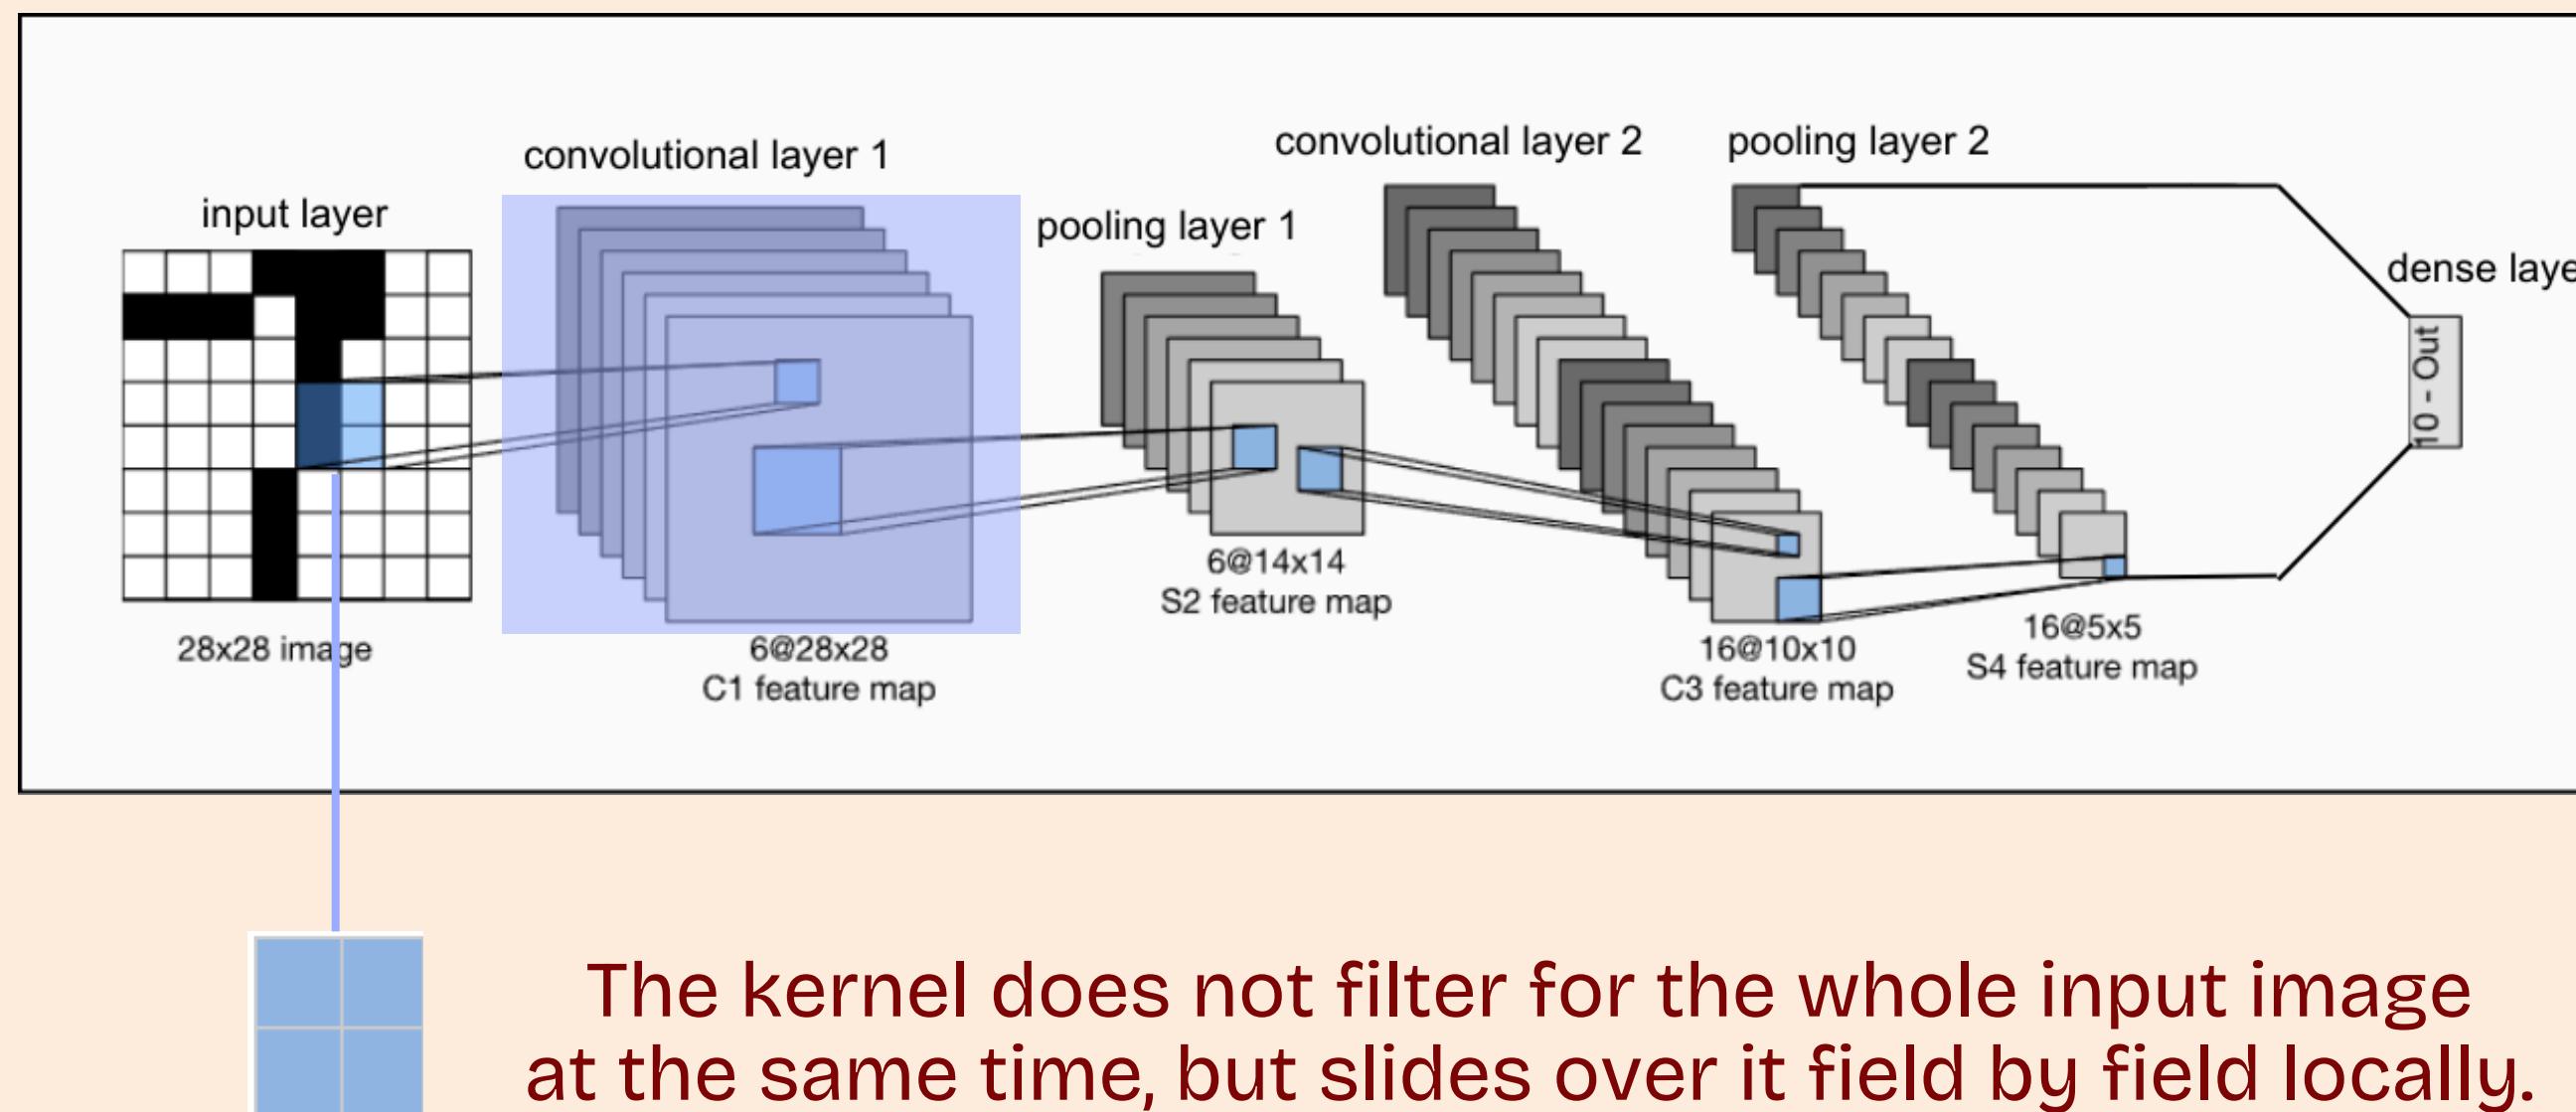
Filter



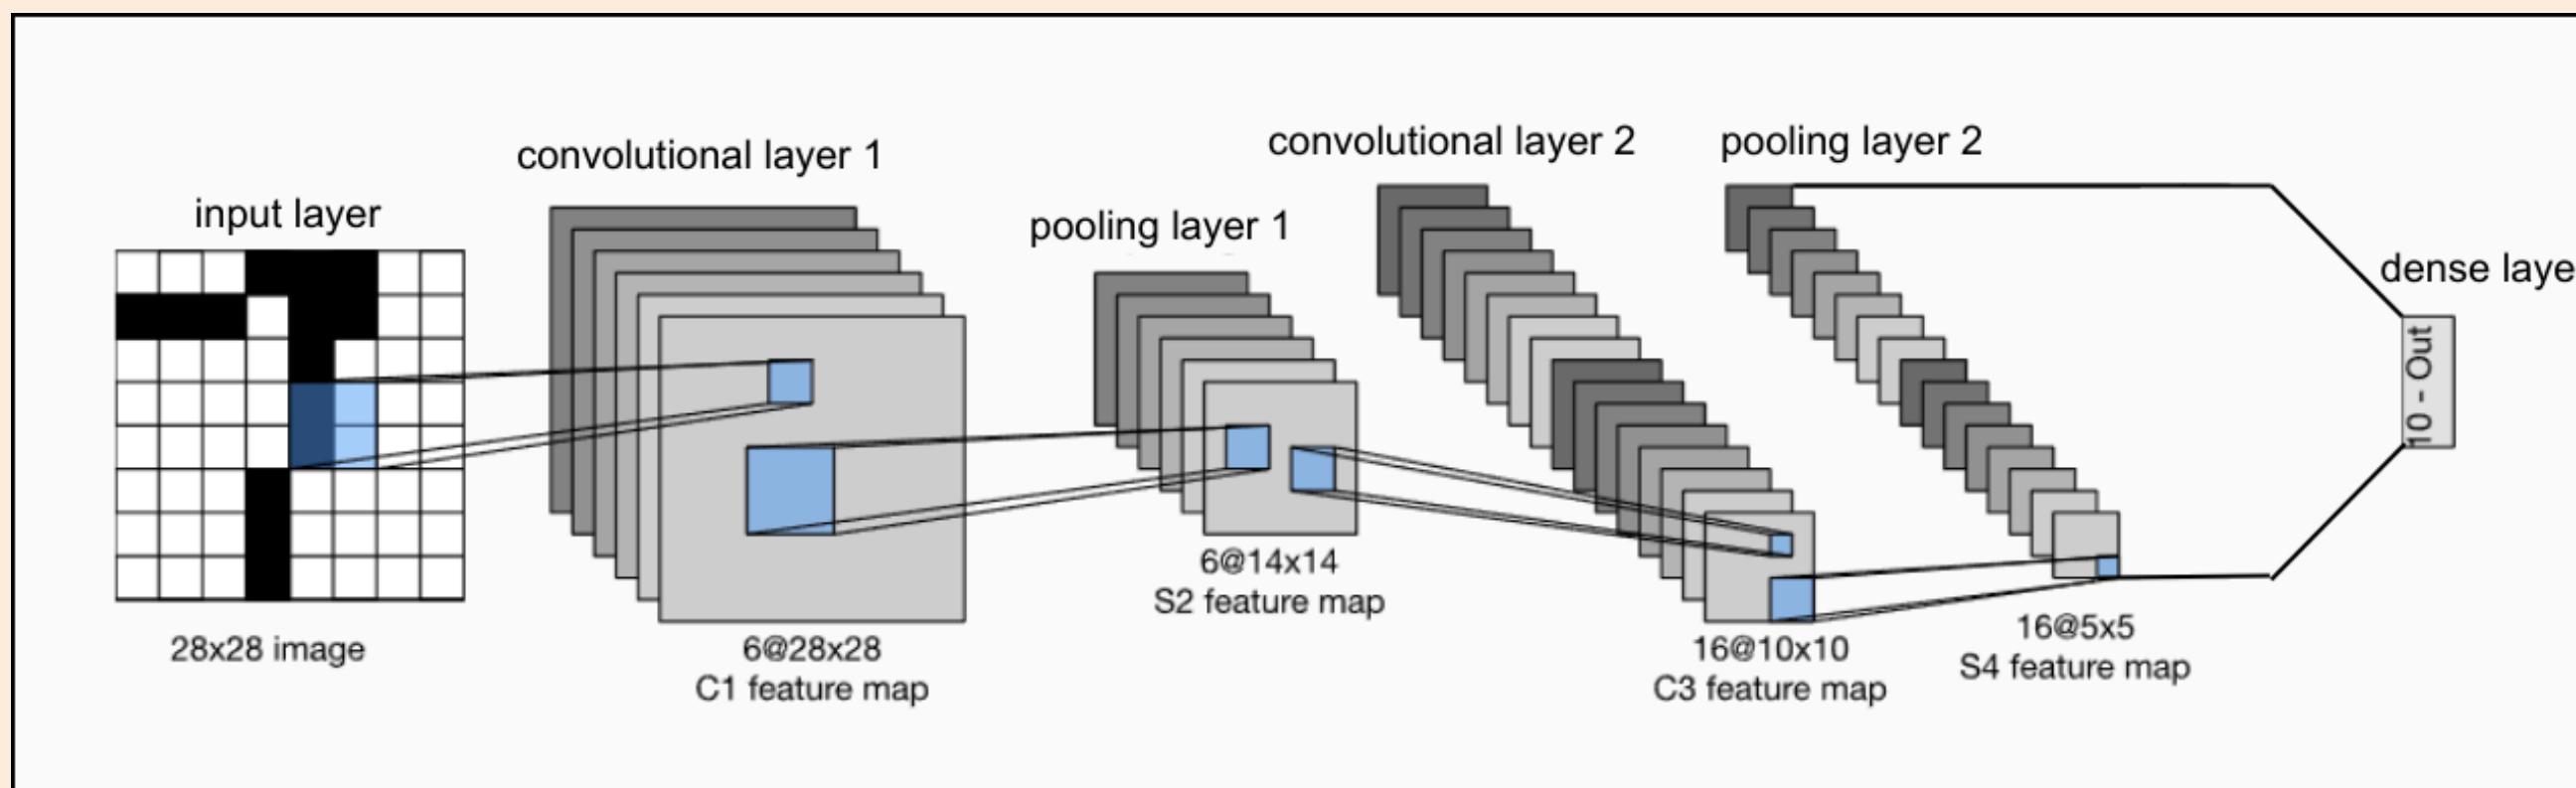
Kernel



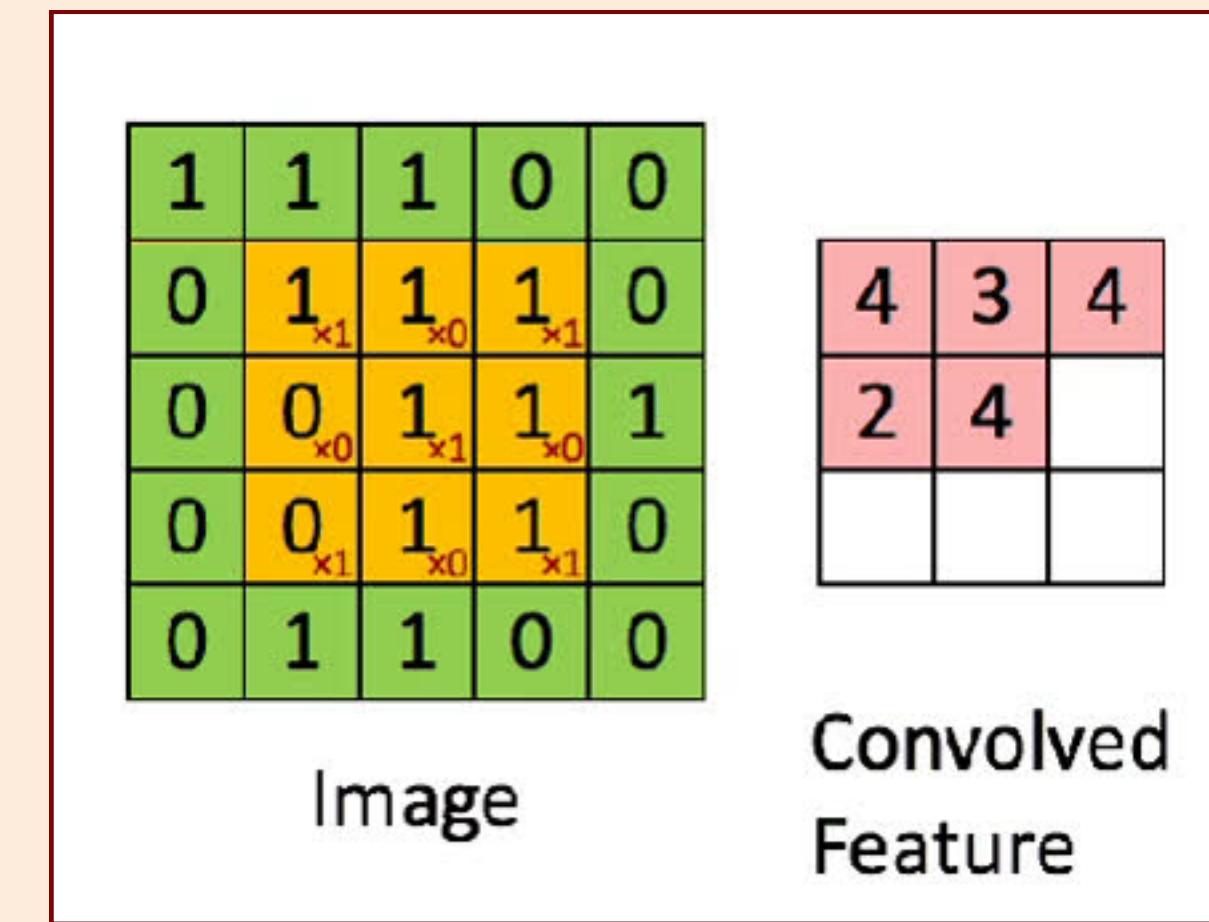
Local receptive field



Convolution

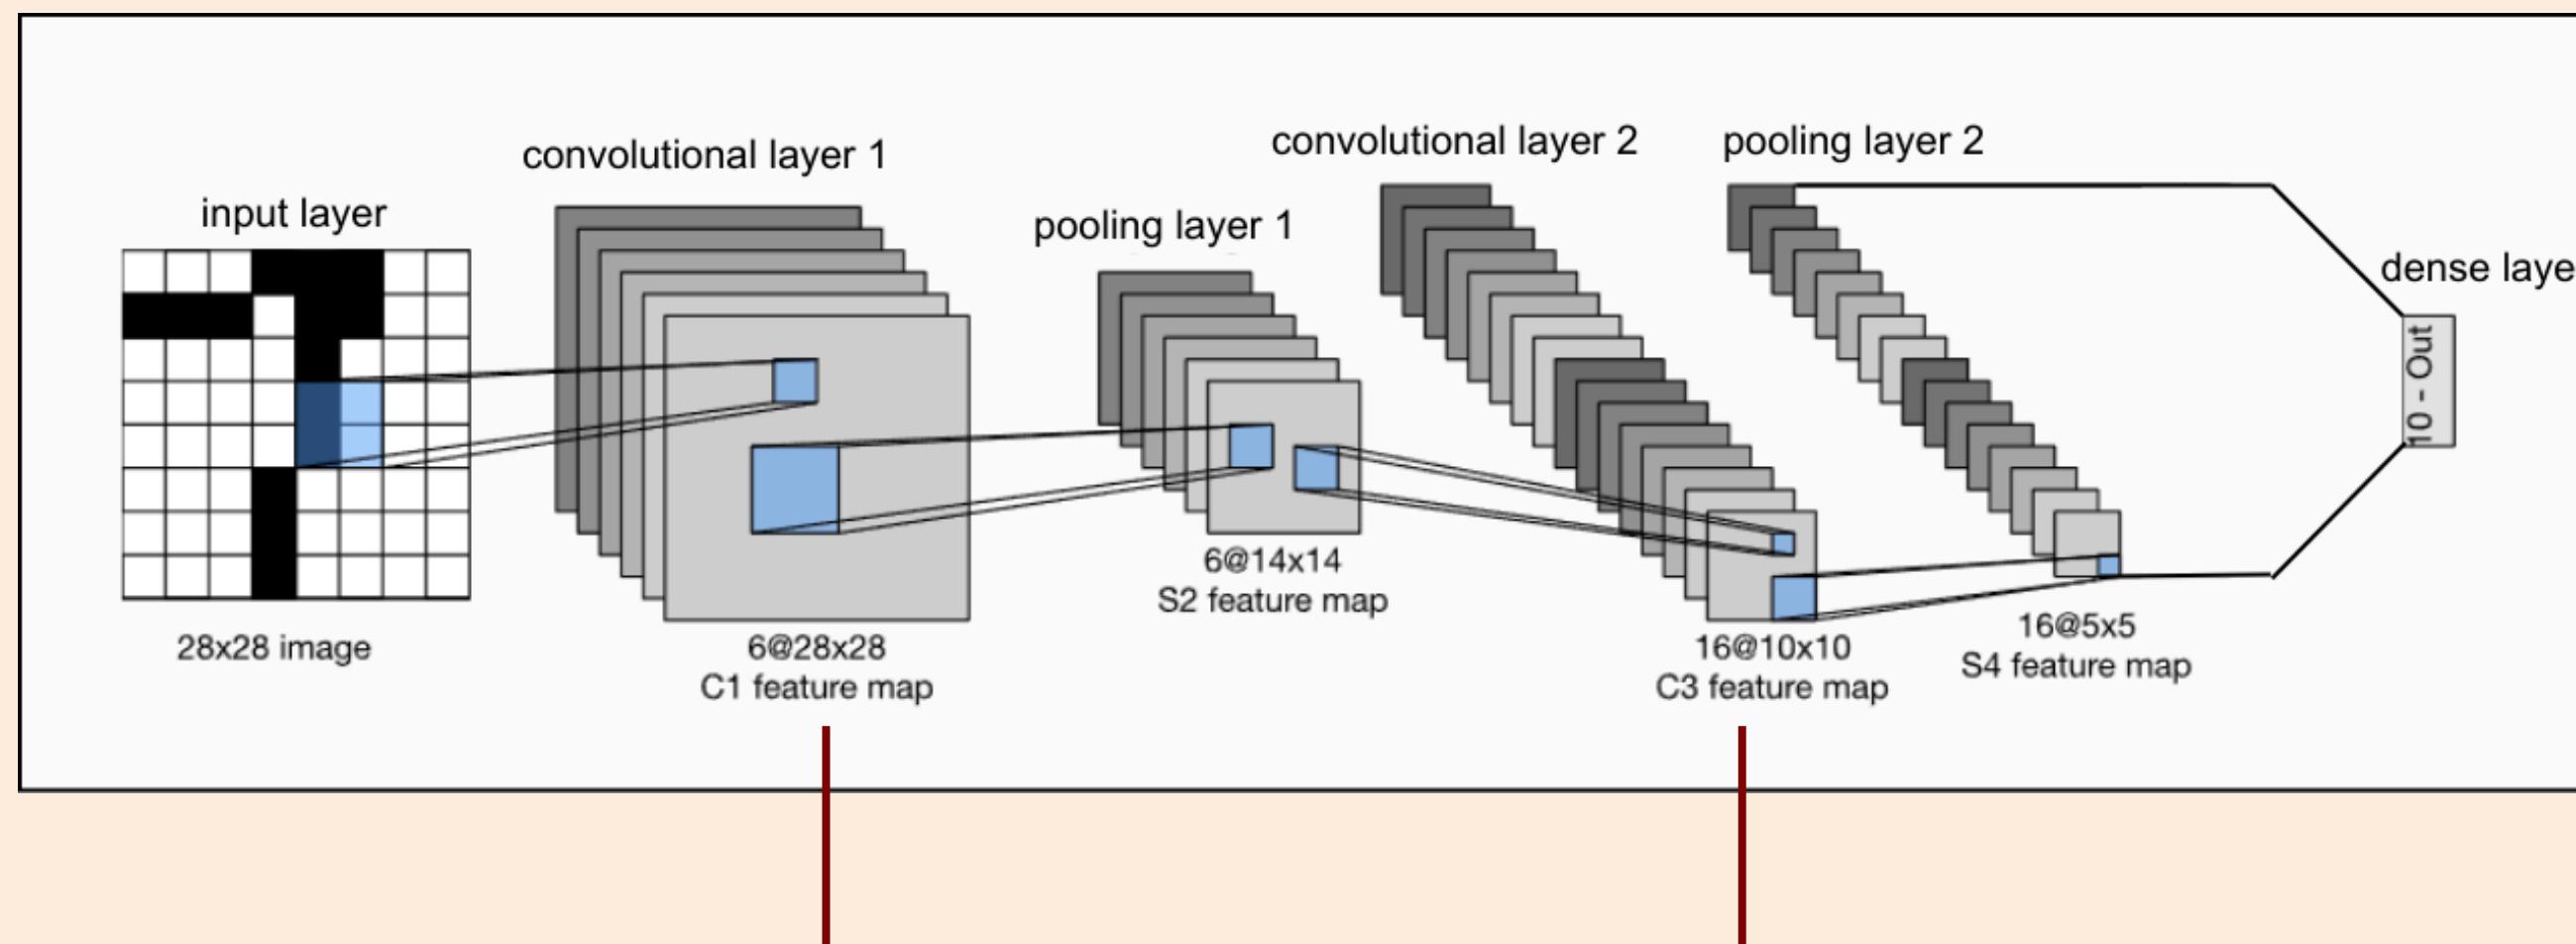


Our computer performs this using an operation called **convolution**.



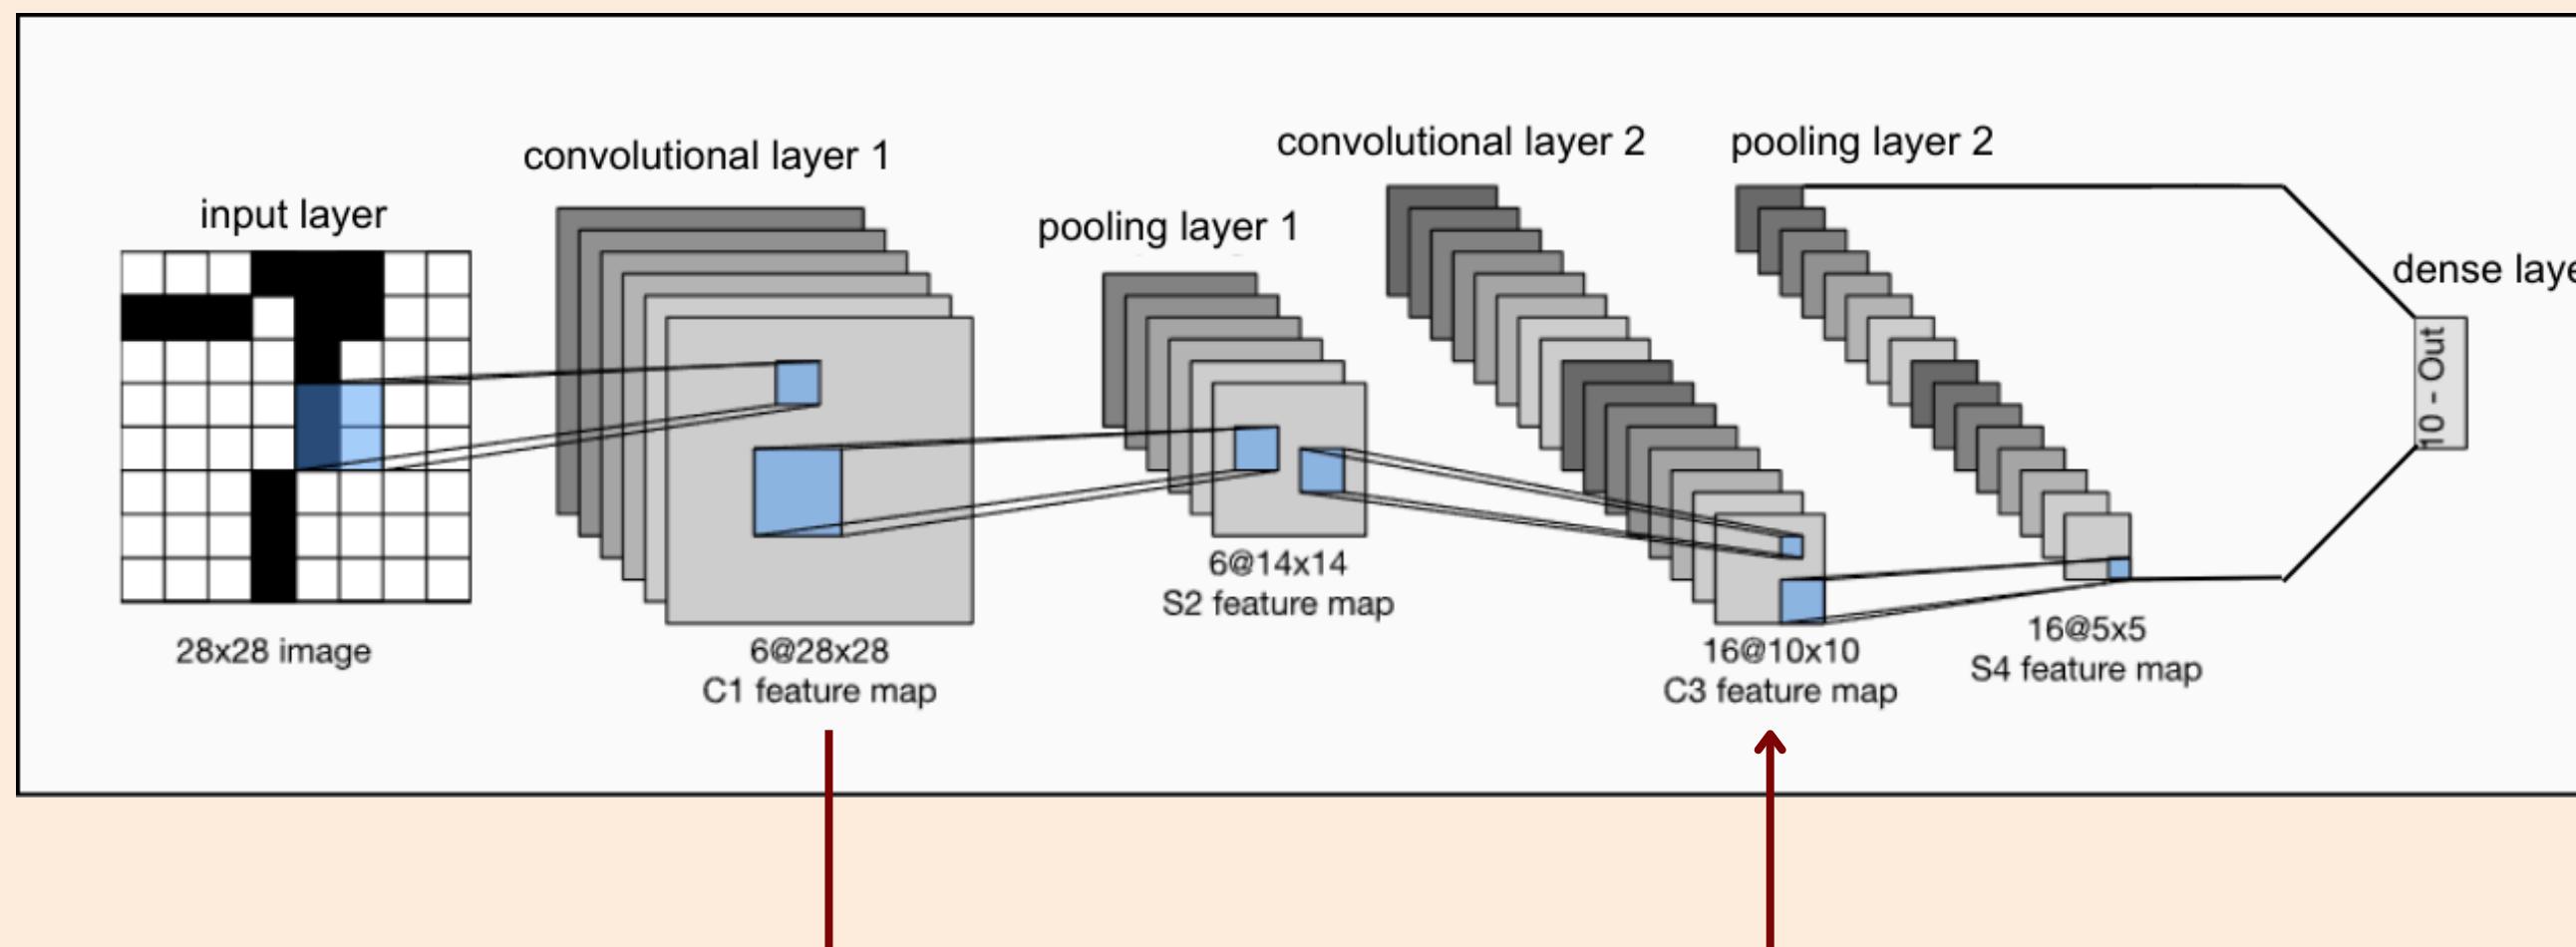
Convolution simply creates an output matrix by **sliding** over the input, **summing** up element-wise **products** between each receptive field and the kernel.

Convolutional layer



Therefore, these layers are called **convolutional layers**.

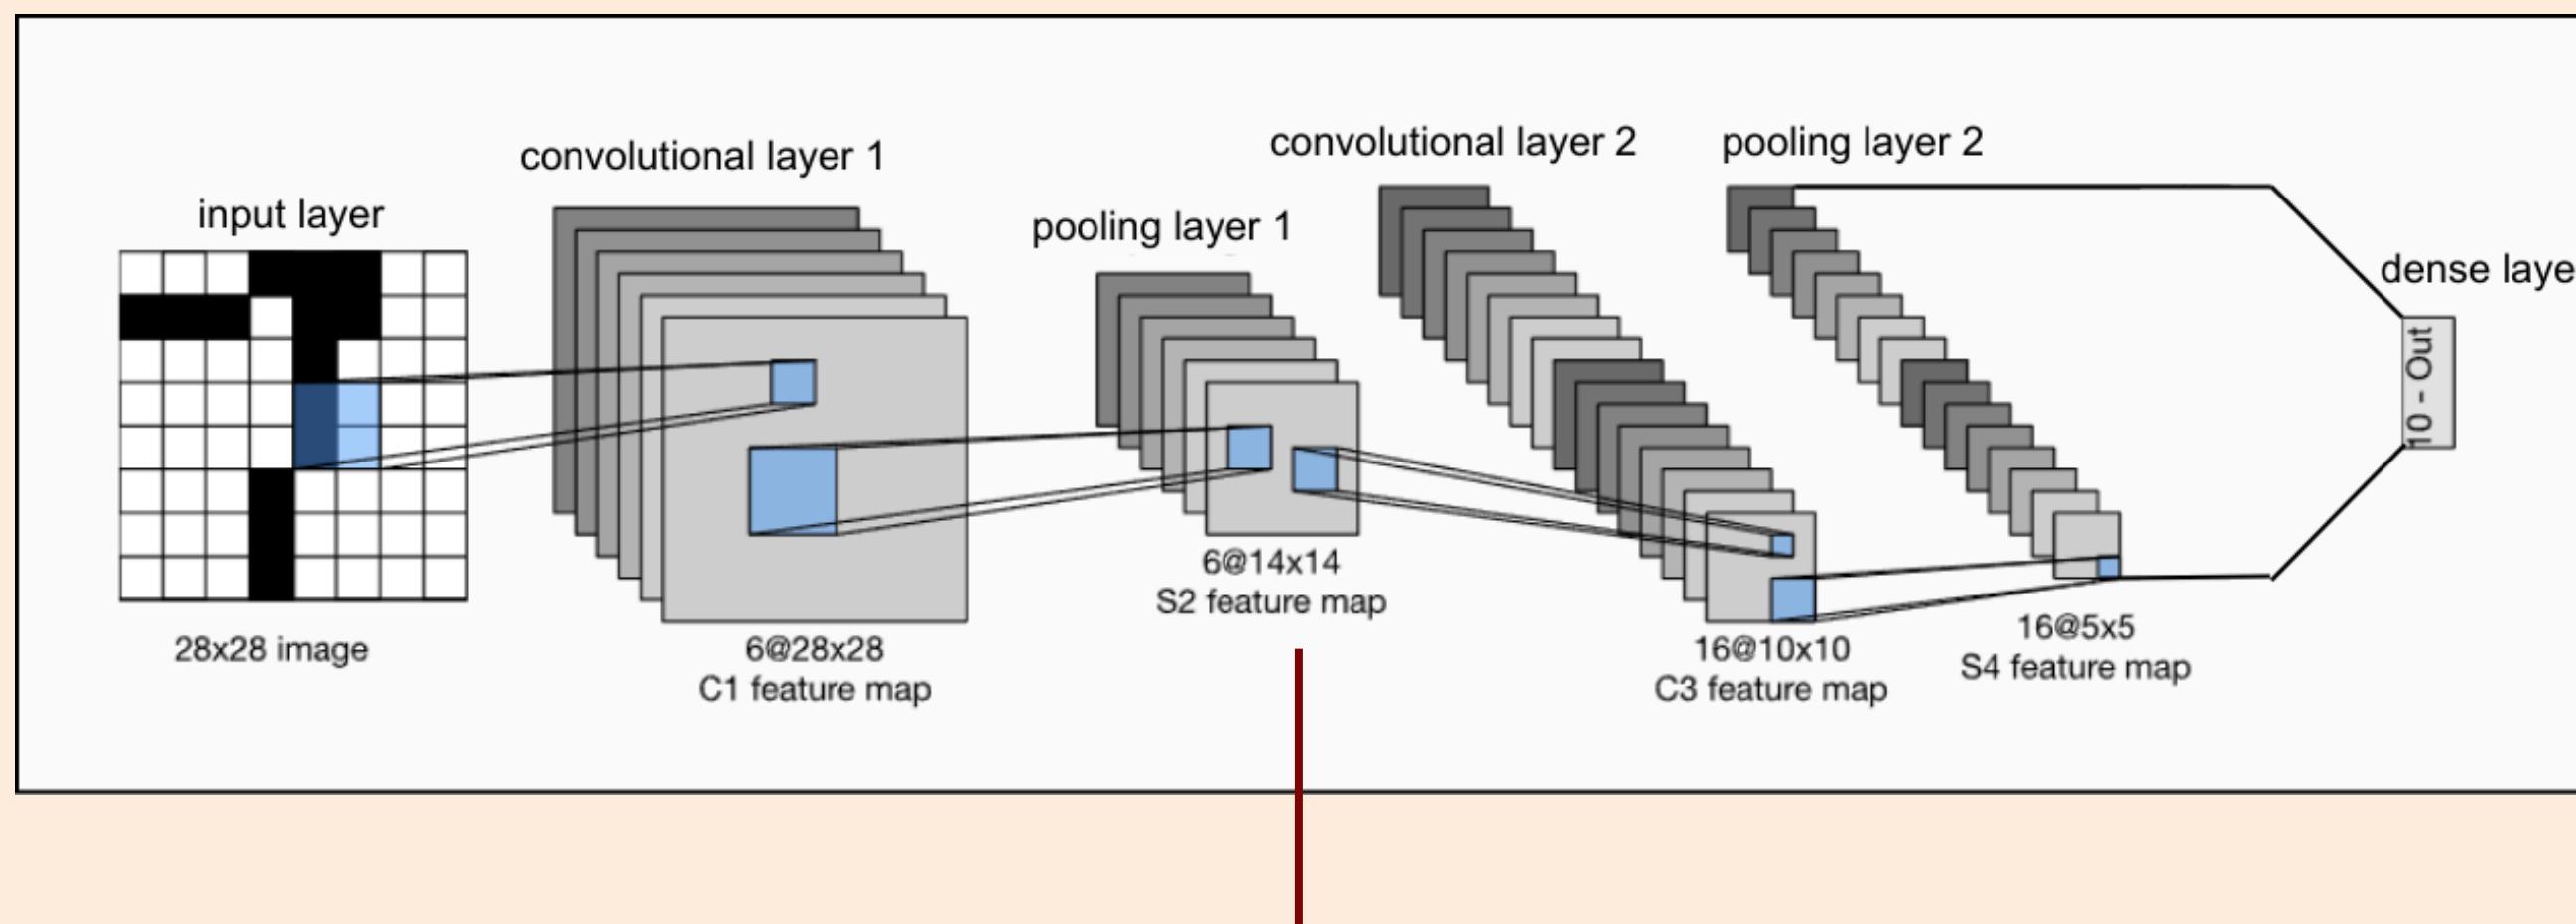
Downsampling



We may want to **reduce the size*** of outputted feature maps before inputting them into deeper layers.

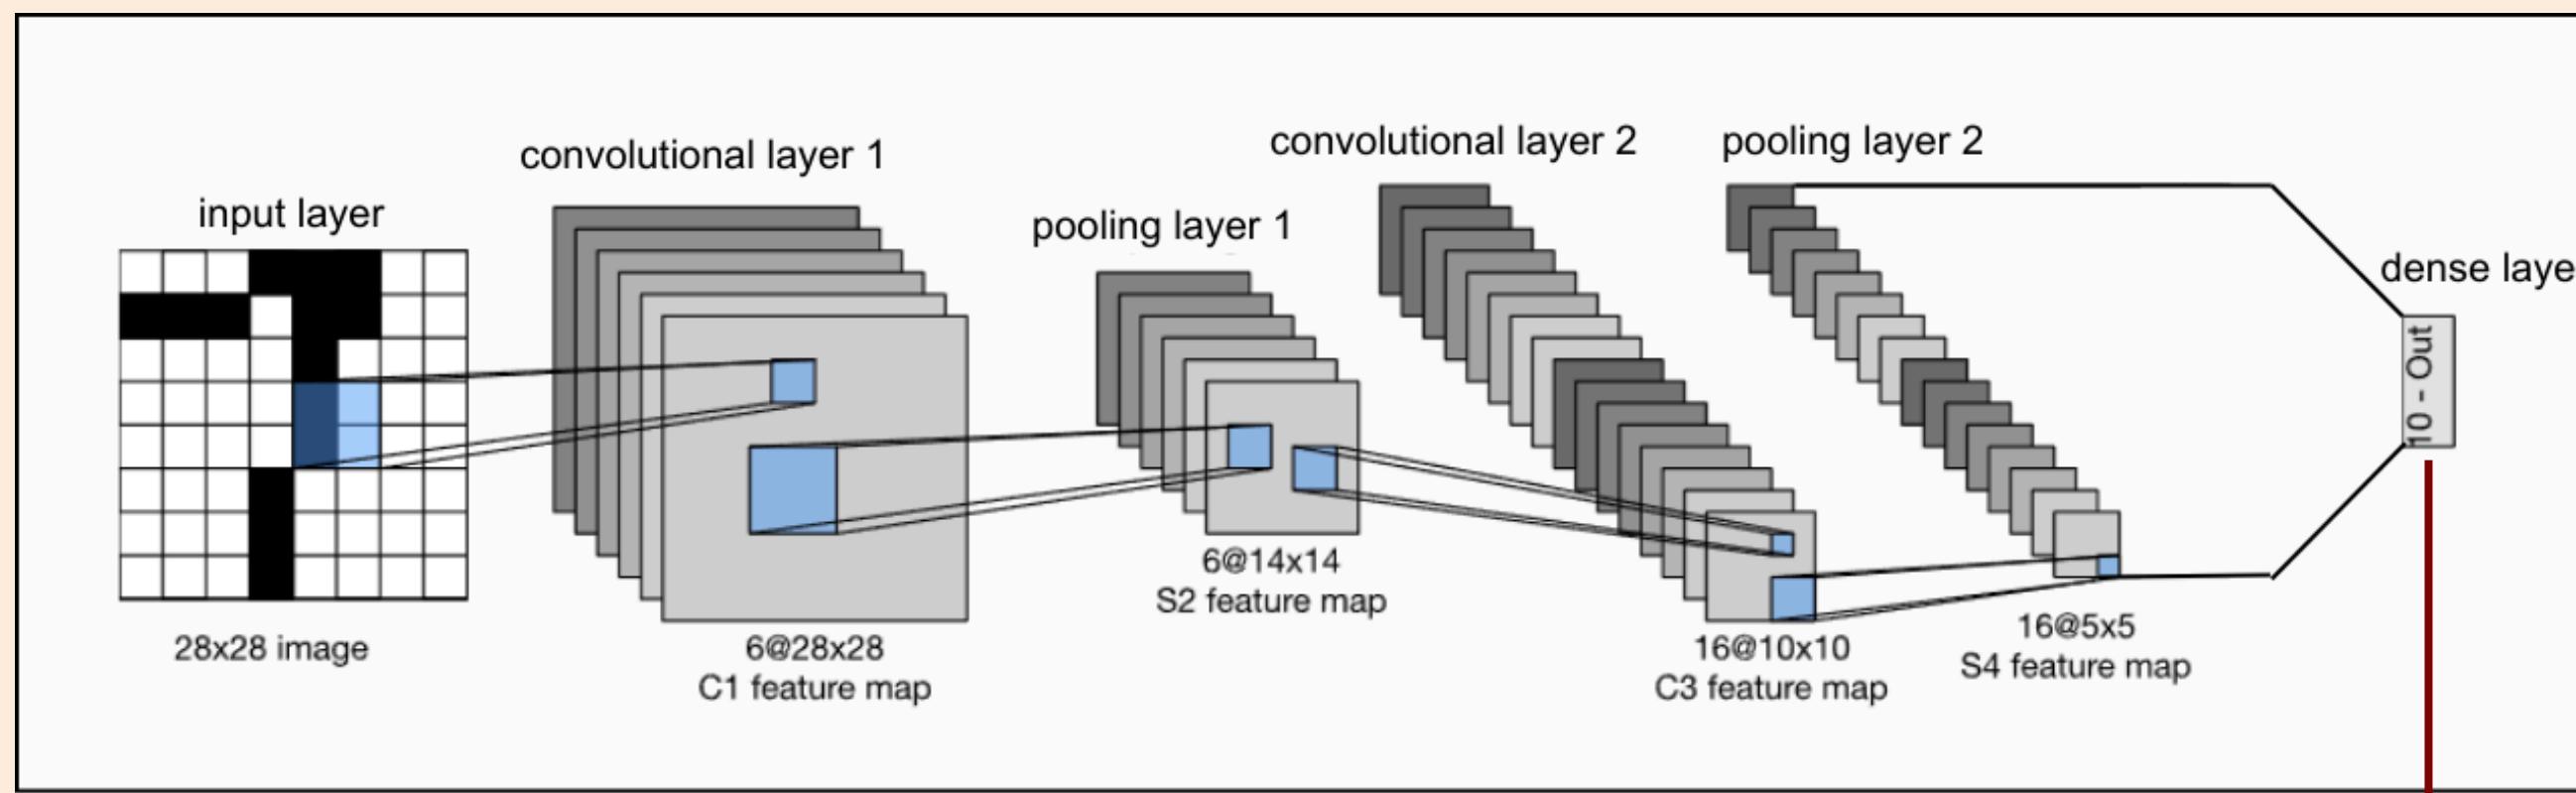
* to reduce computational cost, sensitivity to noise and overfitting

Pooling layer



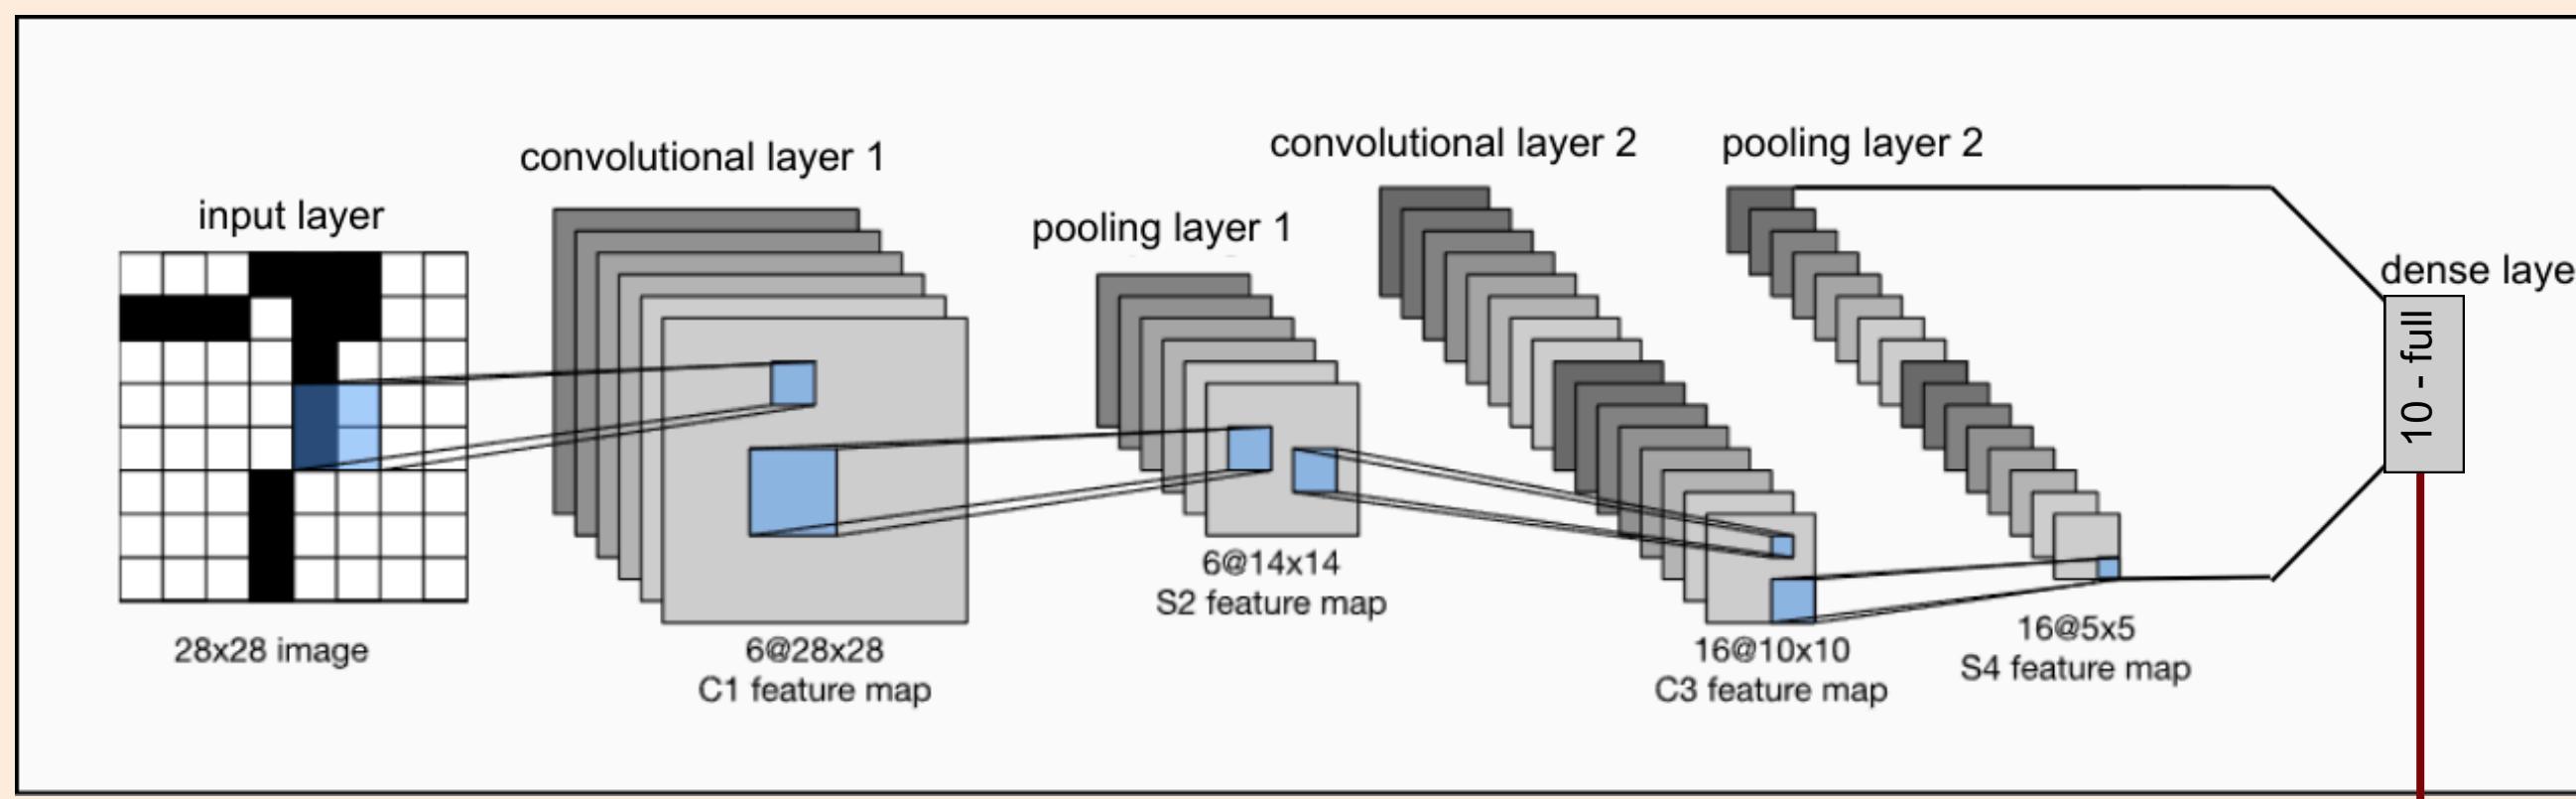
Therefore, we use an intermediate layer called **pooling layer**.
Pooling simply means reducing the size!

Dense layer



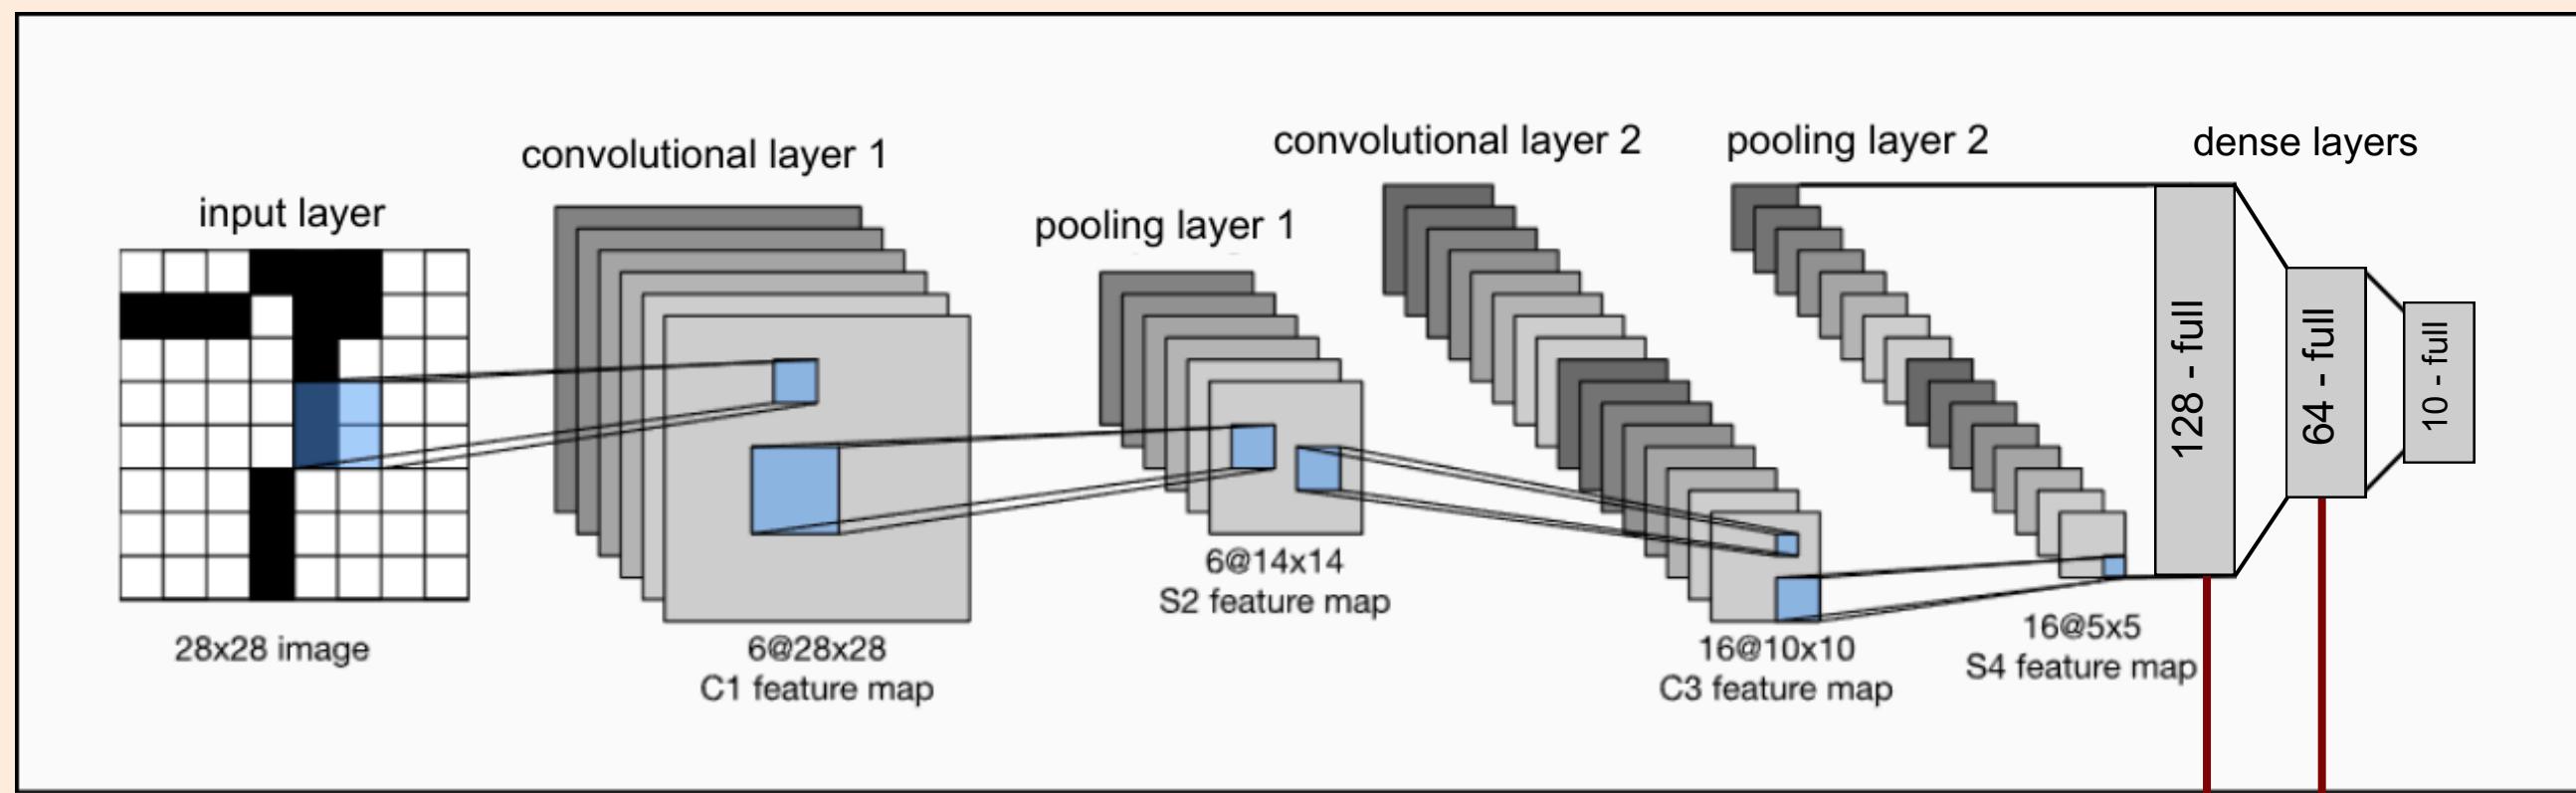
At the end of the network, we perform reasoning based on **all** features to do the classification.

Dense layer



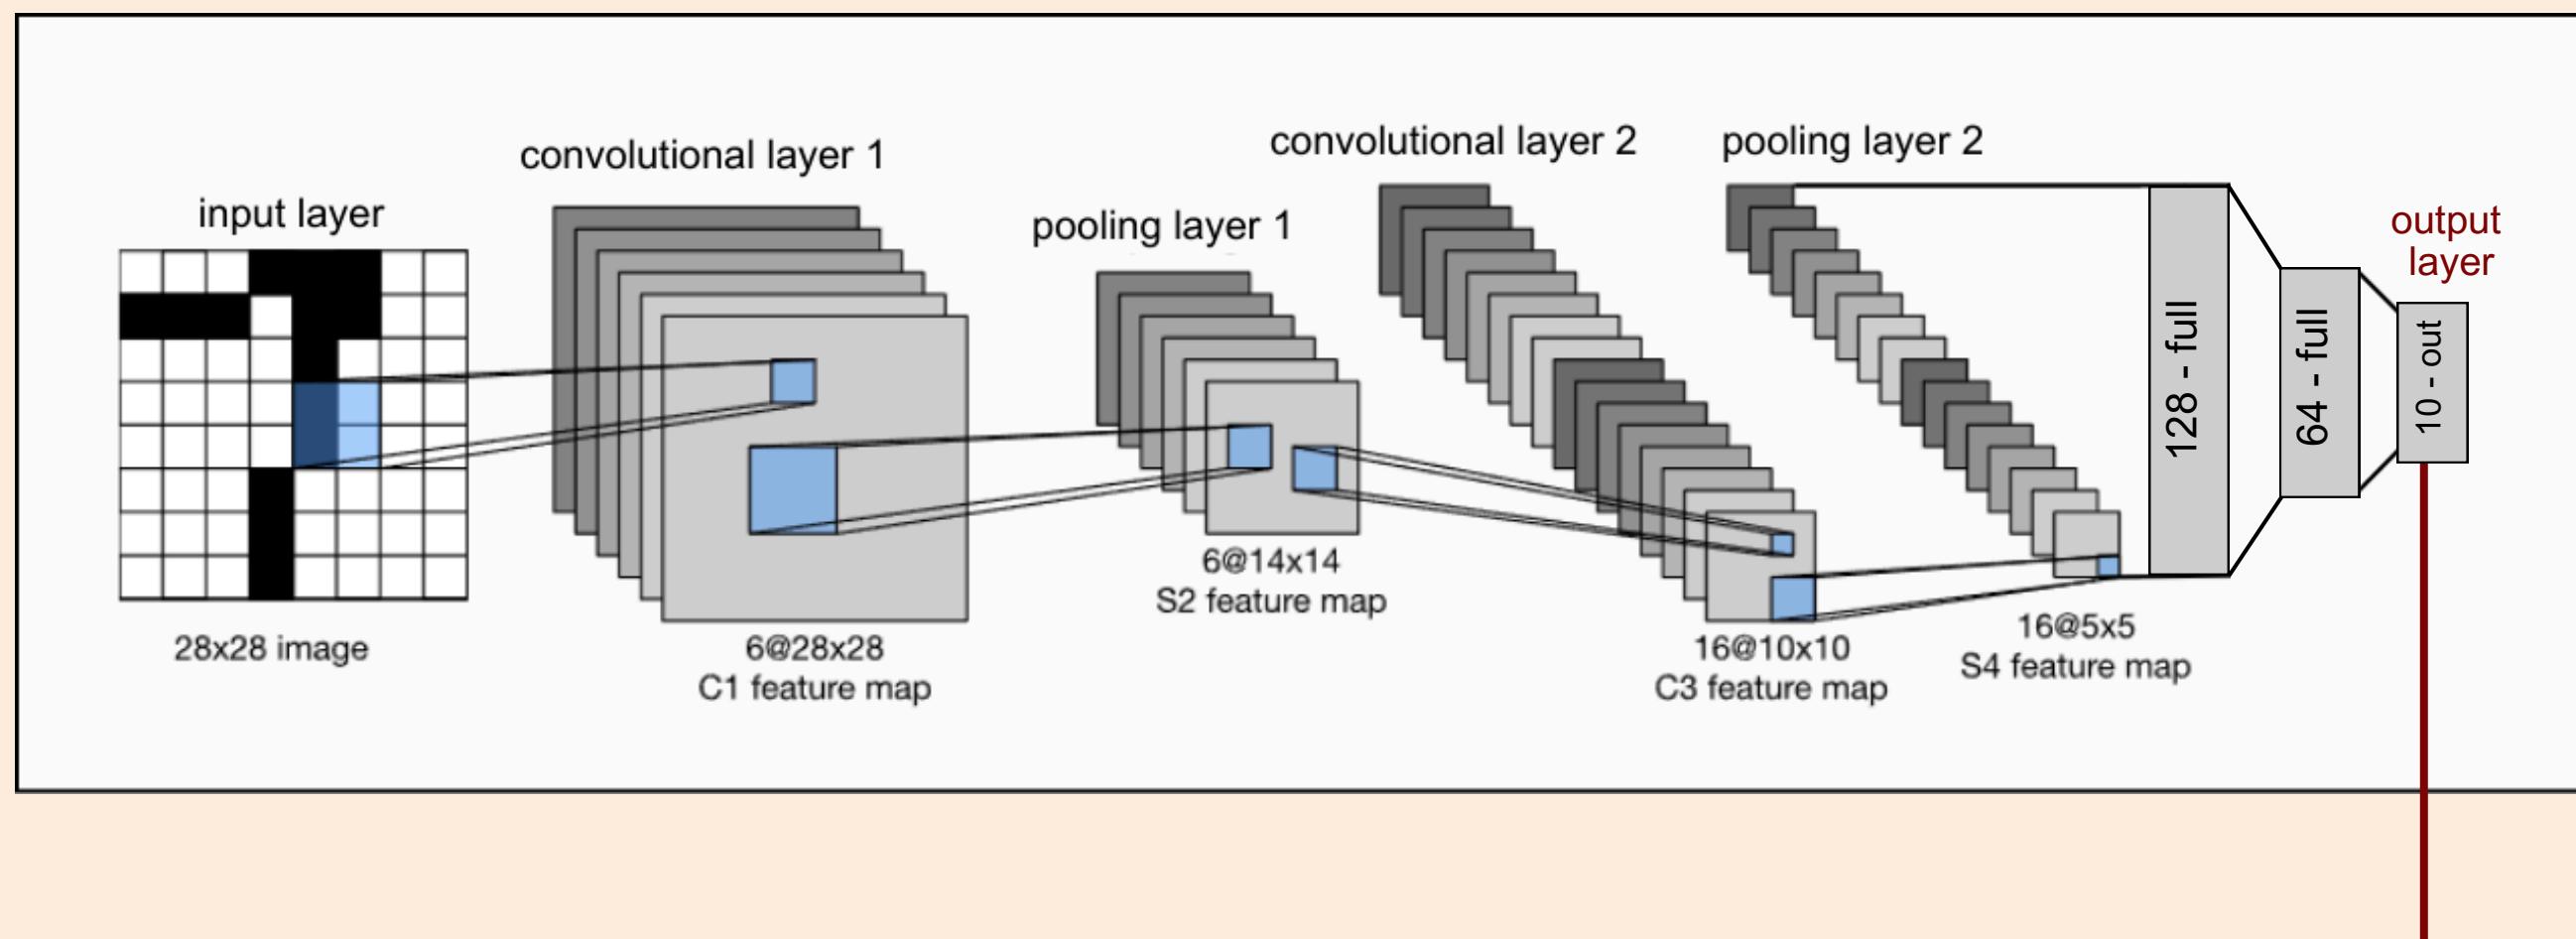
We **connect** each neuron in this layer to **all neurons** of the previous layer. Therefore, we call it a **fully connected** layer, or **dense** layer.

Dense layer



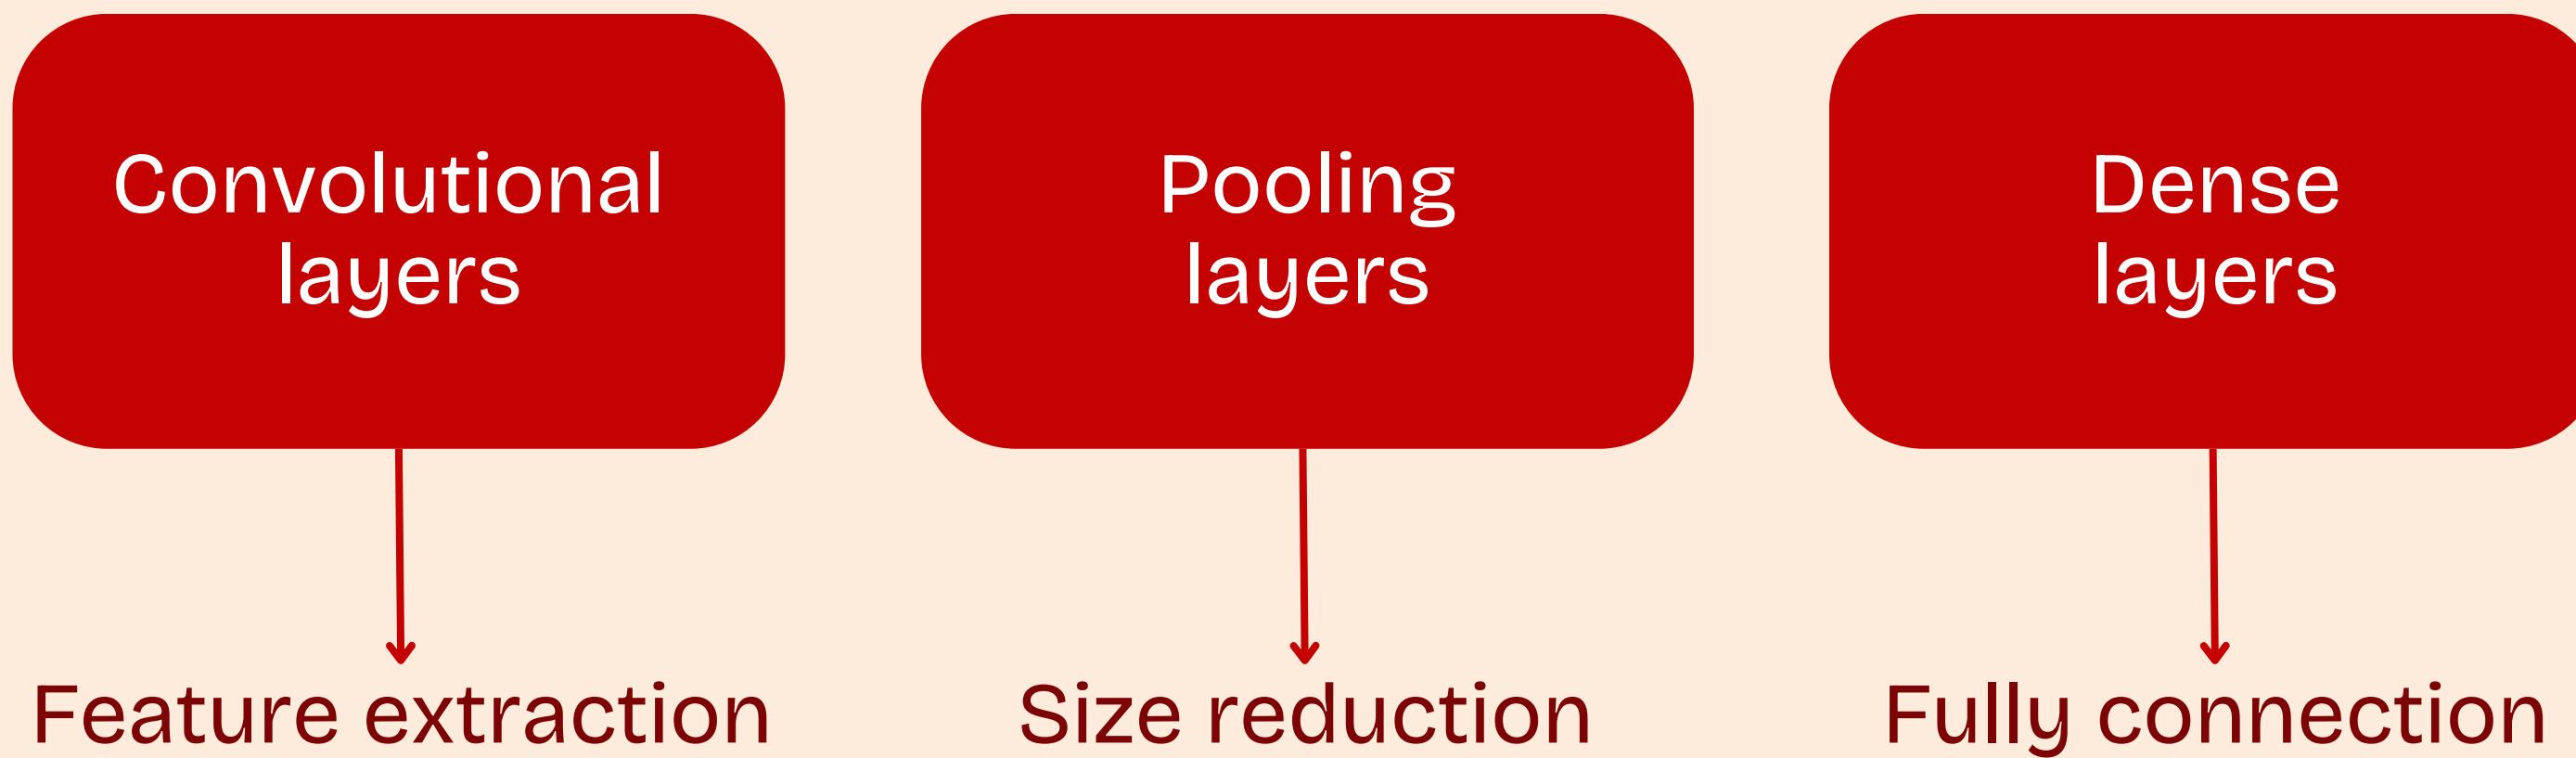
We can add multiple **intermediate dense layers** to make the final classification decision gradually.

Output layer



The final dense layer is called **output layer**.

The idea of our CNN model
is very simple and straightforward!



```
max_epoch = 10
batch_size = 32
tf.reset_default_graph()
input = tf.placeholder(dtype=tf.float32, shape=(None,28,28,1), name='input')
output = tf.placeholder(dtype=tf.float32, shape=(None,10), name='output')

def train_neural_network(x):

    prediction = convolutional_neural_network(x)
    cost = CrossEntropy(prediction, output)
    optimizer = tf.train.AdamOptimizer(learning_rate=0.01).minimize(cost)
    correct = tf.equal(tf.argmax(prediction, 1), tf.argmax(output, 1))
    accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))

    with tf.Session() as sess:
        sess.run(tf.global_variables_initializer())
        indices = np.arange(len(X_train))

        for epoch in range(max_epoch):

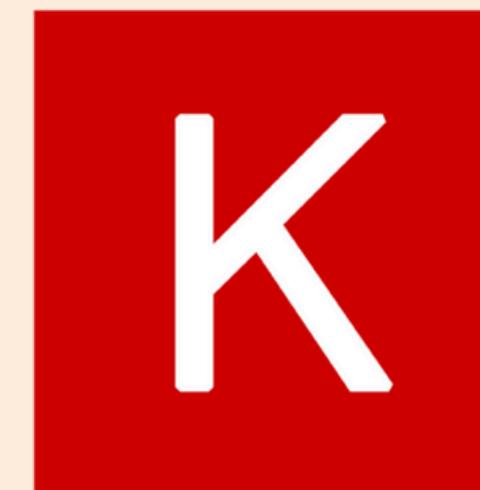
            for i in range(0, len(X_train), batch_size):

                epoch_x = X_train[indices[i : i+batch_size]]
                epoch_y = y_train[indices[i : i+batch_size]]
                a,b,c = sess.run([accuracy, optimizer, cost], feed_dict={input: epoch_x, output:
                    epoch_y})

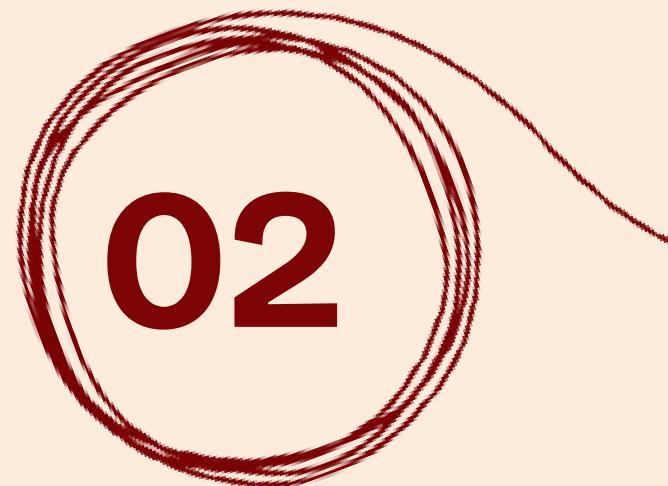
            result = sess.run([cost, accuracy],feed_dict={input:X_test, output:y_test})
```

But it's quite complicated
to implement, especially
for beginners like us.

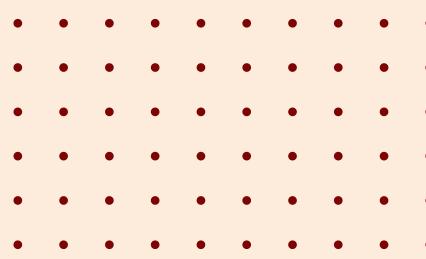
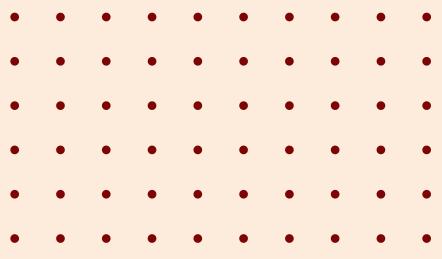
Don't worry. Keras will help you!



Keras



02 Keras





Keras was created by Francois Chollet,
an AI researcher and software engineer, in **2015**.

At that time, Theano was a common framework for DL.

theano

However, it's a **low-level** framework.

Working with low-level framework seems,
most of the time, complicated and time-consuming!

Caffe

He wanted a high-level tool, and Caffe was an option at that time.

```
layer {
    name: "conv1"
    type: "Convolution"
    bottom: "data"
    top: "conv1"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 20
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "xavier"
        }
        bias_filler {
            type: "constant"
        }
    }
}
```

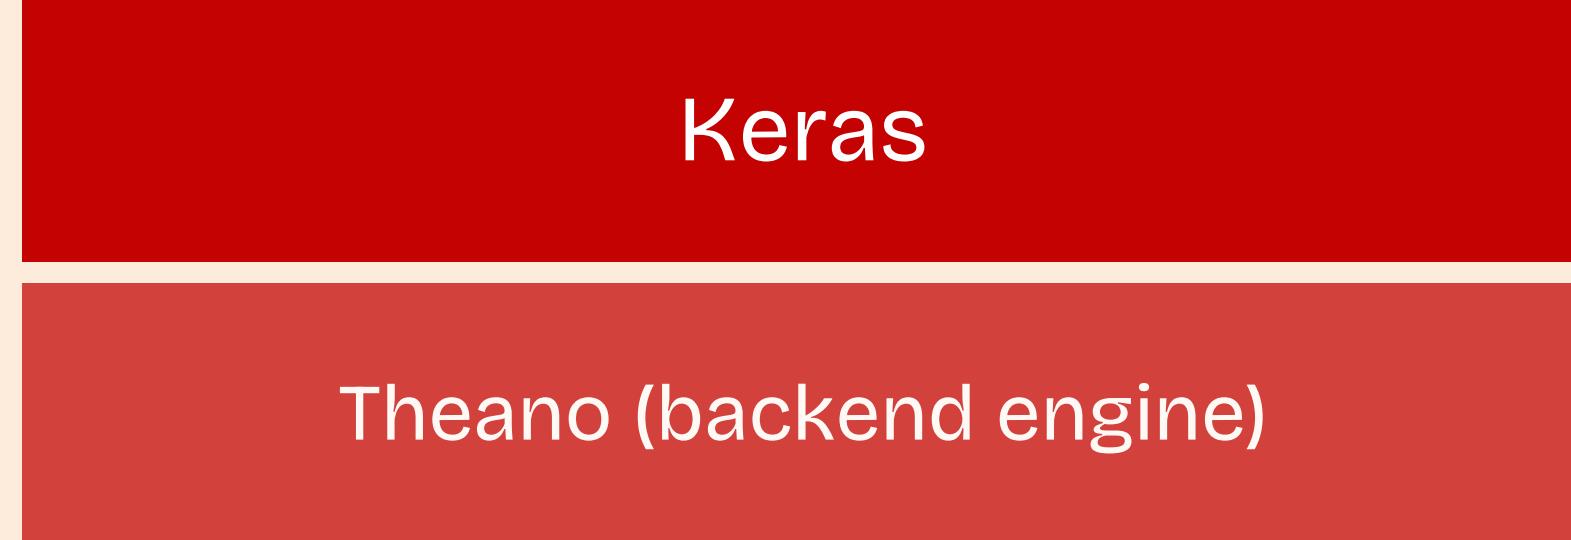
```
layer {
    name: "pool1"
    type: "Pooling"
    bottom: "conv1"
    top: "pool1"
    pooling_param {
        pool: MAX
        kernel_size: 2
        stride: 2
    }
}
```

```
layer {
    name: "prob"
    type: "Softmax"
    bottom: "ip2"
    top: "prob"
}
```

In Caffe, we have to define the model architecture in a configuration file called **prototxt**.

It is **inconvenient** and **inflexible**.

And this motivated Francois to create an
easy-to-use and **flexible** high-level deep learning framework...



Keras

A diagram illustrating the relationship between Keras and Theano. It consists of two stacked horizontal bars. The top bar is dark red and contains the word "Keras" in white. The bottom bar is a lighter shade of red and contains the text "Theano (backend engine)" in white.

Theano (backend engine)

... on top of Theano as backend engine, in 2015.

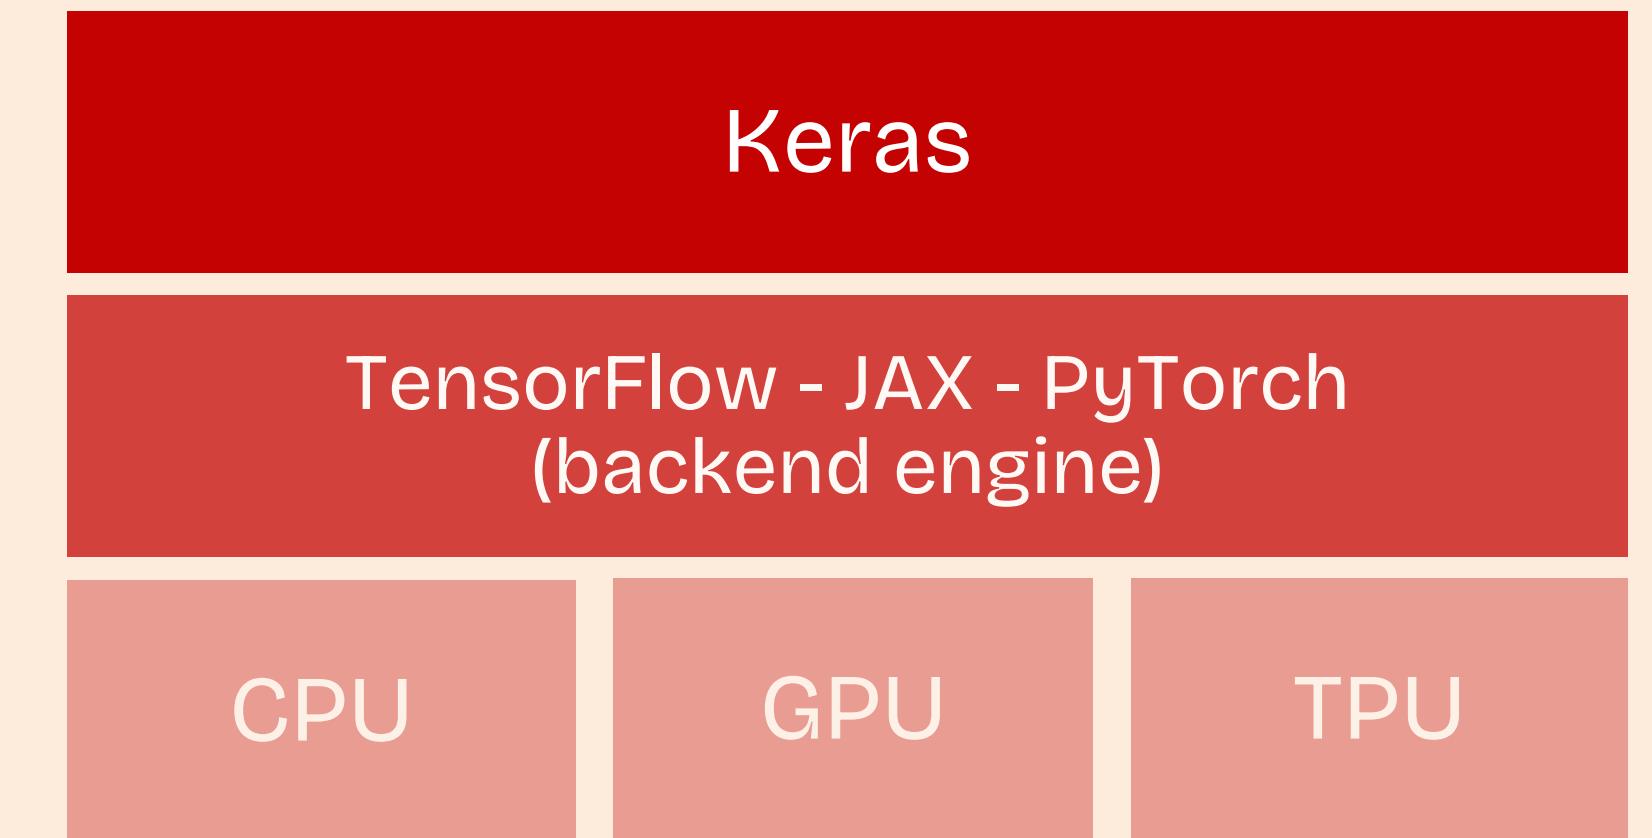
In 2023, Keras 3 was released as a multi and cross-backend DL framework.

Keras

TensorFlow - JAX - PyTorch

* Note: since version 2, Keras has also been integrated into core TensorFlow.

Keras is a
HIGH-LEVEL
CROSS-BACKEND
deep learning framework.



HIGH LEVEL

Keras abstracts deep learning **Layers** and **Models**.

01 Problem

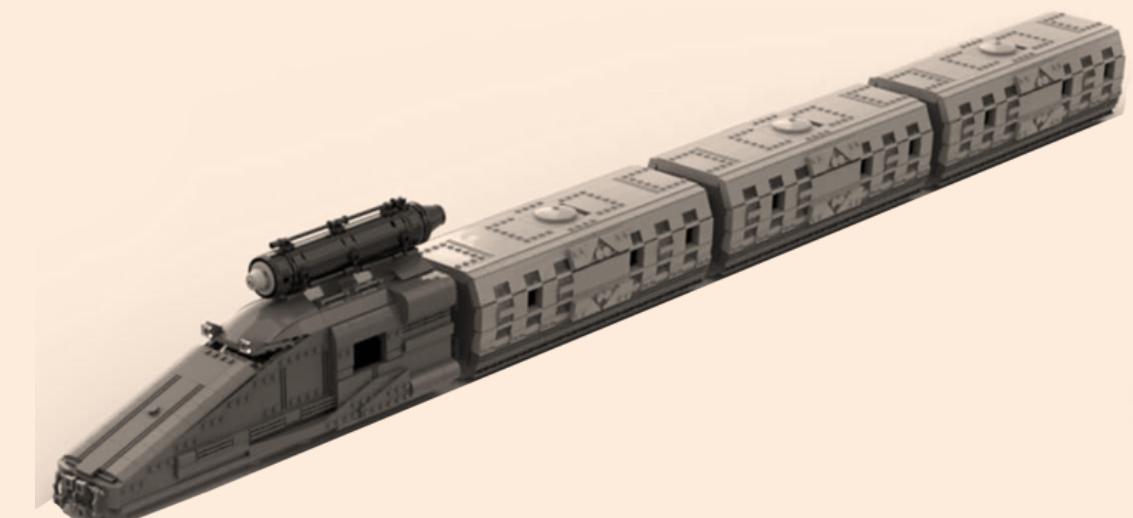


02 Keras

03 Demonstration

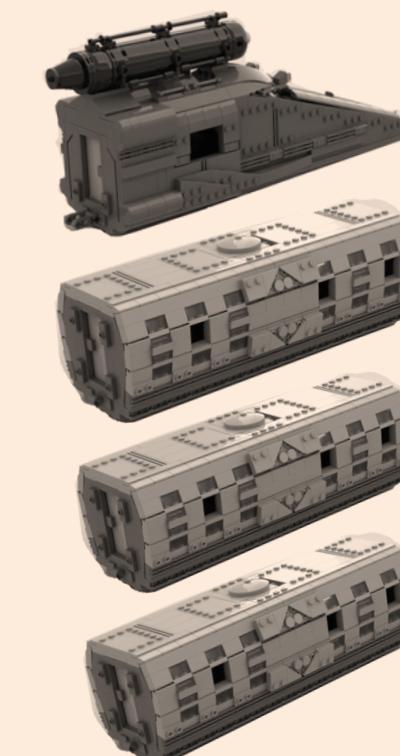
04 Conclusion

Developer



Keras

Developer



```
def conv2d(x, kernel, strides=(1, 1), border_mode='valid',
          dim_ordering='th', image_shape=None, filter_shape=None):

    if border_mode == 'same':
        padding = 'SAME'
    elif border_mode == 'valid':
        padding = 'VALID'
    else:
        # raise exception

    strides = (1,) + strides + (1,)

    if _FLOATX == 'float64':
        x = tf.cast(x, 'float32')
        kernel = tf.cast(kernel, 'float32')

    if dim_ordering == 'th':
        x = tf.transpose(x, (0, 2, 3, 1))
        kernel = tf.transpose(kernel, (2, 3, 1, 0))
        x = tf.nn.conv2d(x, kernel, strides, padding=padding)
        x = tf.transpose(x, (0, 3, 1, 2))
    elif dim_ordering == 'tf':
        x = tf.nn.conv2d(x, kernel, strides, padding=padding)
    else:
        raise Exception('Unknown dim_ordering: ' + str(dim_ordering))

    if _FLOATX == 'float64':
        x = tf.cast(x, 'float64')

    return x
```

Keras
abstracts layers

```
Conv2D(32, kernel_size=(3,3),  
       activation='relu',  
       input_shape=(28,28,1),  
       padding='SAME')
```

```

class Sequential(Model):
    def __init__(self, layers=[], name=None):
        ...
    def add(self, layer):
        ...
    def call(self, x, mask=None):
        ...
    def build(self, input_shape=None):
        ...
    def get_weights(self):
        ...
    def set_weights(self, weights):
        ...

    @property
    def validation_data(self):
        ...
    @property
    def training_data(self):
        ...

    def compile(self, optimizer, loss, metrics=[],
               sample_weight_mode=None, **kwargs):
        ...
    def fit(self, x, y, batch_size=32, nb_epoch=10, verbose=1, callbacks=[],
            validation_split=0., validation_data=None, shuffle=True,
            class_weight=None, sample_weight=None, **kwargs):
        ...
    def evaluate(self, x, y, batch_size=32, verbose=1,
                sample_weight=None, **kwargs):
        ...
    def predict(self, x, batch_size=32, verbose=0):
        ...
    def predict_on_batch(self, x):
        ...
    def test_on_batch(self, x, y, sample_weight=None, **kwargs):
        ...

```

```

def add(self, layer):
    if not self.outputs:
        if len(layer.inbound_nodes) == 0:
            # create an input layer
            if not hasattr(layer, 'batch_input_shape'):
                # raise Exception
                batch_input_shape = layer.batch_input_shape
            if hasattr(layer, 'input_dtype'):
                input_dtype = layer.input_dtype
            else:
                input_dtype = None
            layer.create_input_layer(batch_input_shape, input_dtype)

        if len(layer.inbound_nodes) != 1:
            # raise Exception

        if len(layer.inbound_nodes[0].output_tensors) != 1:
            # raise Exception

        self.outputs = [layer.inbound_nodes[0].output_tensors[0]]
        self.inputs = get_source_inputs(self.outputs[0])

        # create an input node
    else:
        # add output layer

    self.layers.append(layer)
    self.built = False

```

Keras abstracts models

```

model = Sequential()
model.add(Conv2D(...))
model.add(MaxPooling2D(...))
model.add(Conv2D(...))
...
model.add(Dense(...))

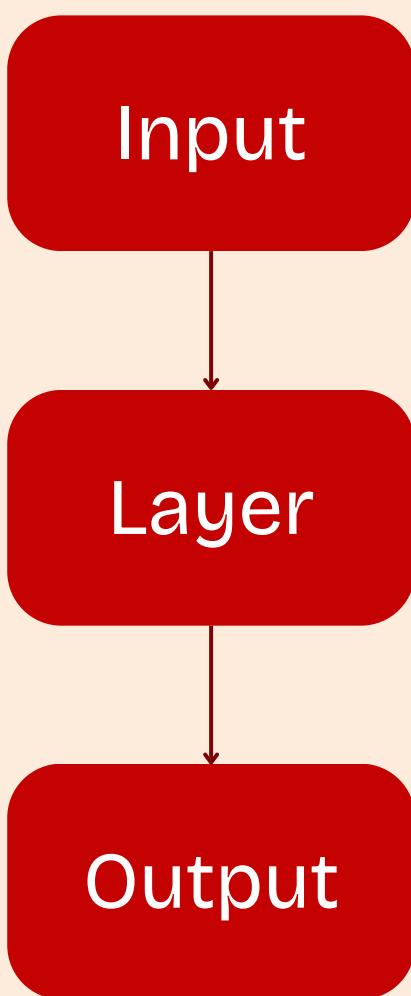
model.compile(loss=..., optimizer=..., metrics=...)

history_data = model.fit(...)

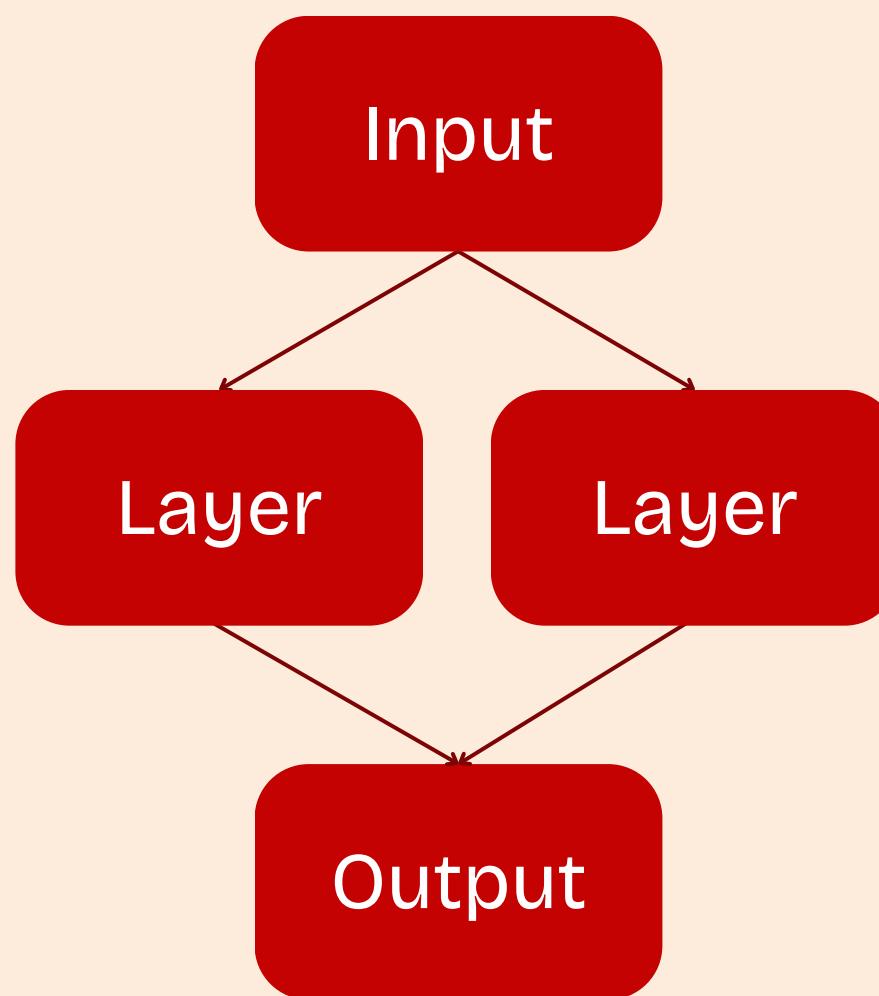
loss,accuracy = model.evaluate(...)

```

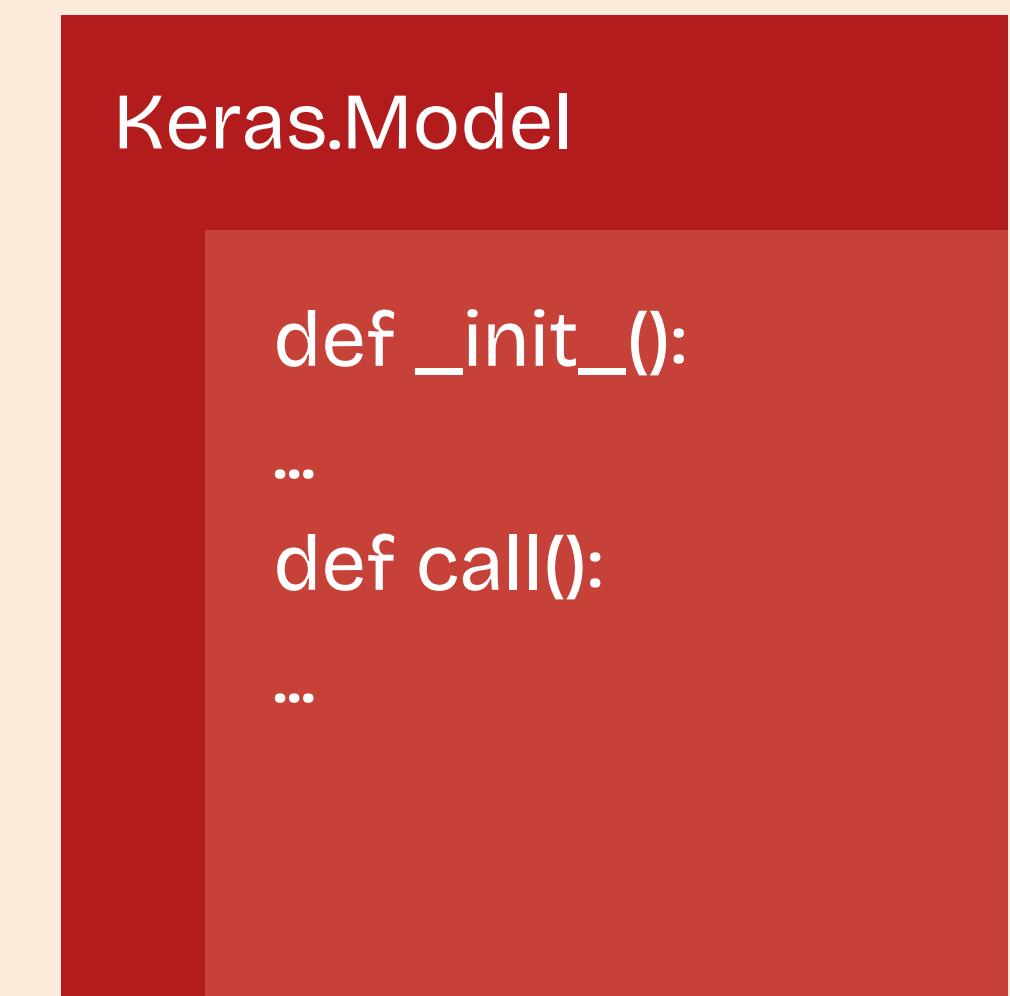
Keras abstracts 3 types of models



Sequential

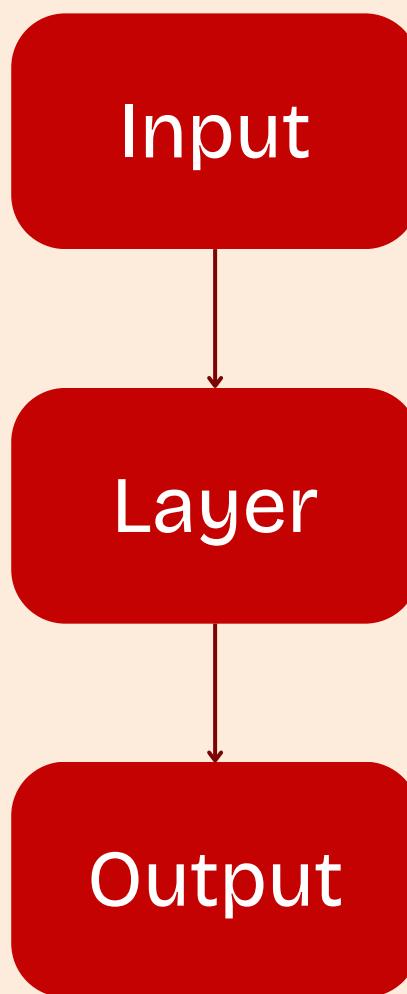


Functional API



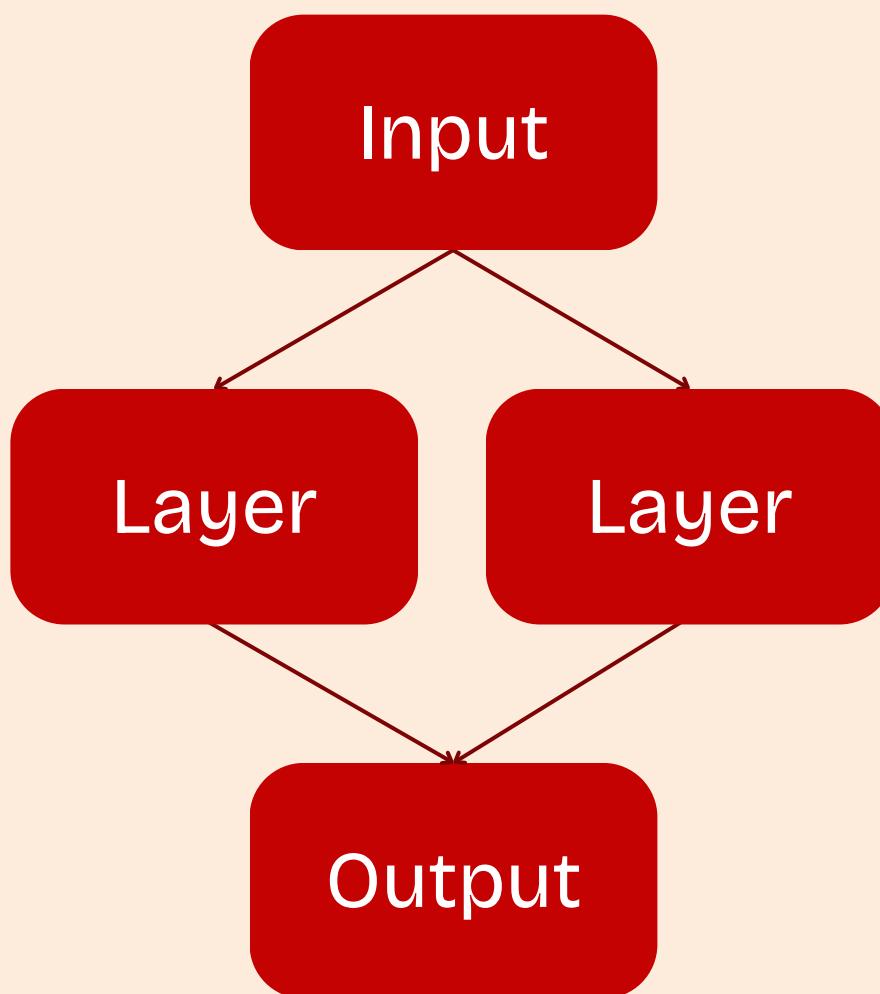
Model Subclassing

Sequential Model



```
# Build the model
model = Sequential([
    Flatten(input_shape=(28, 28)),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])
```

Functional API



```
# Define the input
inputs = layers.Input(shape=(28, 28, 1))

# Define two parallel layers
layer1 = layers.Conv2D(32, (3, 3), activation='relu')(inputs)
layer2 = layers.Conv2D(32, (5, 5), activation='relu')(inputs)

# Merge the outputs of the two layers
merged = layers.concatenate([layer1, layer2])

# Add more layers after merging
x = layers.MaxPooling2D((2, 2))(merged)...
outputs = layers.Dense(10, activation='softmax')(x)
model = Model(inputs=inputs, outputs=outputs)
```

Model Subclassing

Keras.Model

```
def __init__():
...
def call():
...
```

```
class MyModel(keras.Model):
    def __init__(self):
        super(MyModel, self).__init__()
        self.flatten = Flatten()
        self.dense1 = Dense(128, activation='relu')
        self.dense2 = Dense(10, activation='softmax')

    def call(self, inputs):
        x = self.flatten(inputs)
        x = self.dense1(x)
        return self.dense2(x)

# Instantiate the model
model = MyModel()
```

We can compile all these models with this simple code!

```
# Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
model.fit(x_train, y_train, epochs=5, batch_size=32, validation_data=(x_test, y_test))

# Evaluate model
model.evaluate(x_test, y_test)
```

CROSS BACKEND

Keras allows working across TensorFlow, PyTorch and JAX.

One codebase for all backends with keras.ops

```
import keras
from keras import ops

class TokenAndPositionEmbedding(keras.Layer):
    def __init__(self, max_length, vocab_size, embed_dim):
        super().__init__()
        self.token_embed = self.add_weight(
            shape=(vocab_size, embed_dim),
            initializer="random_uniform",
            trainable=True,
        )
        self.position_embed = self.add_weight(
            shape=(max_length, embed_dim),
            initializer="random_uniform",
            trainable=True,
        )

    def call(self, token_ids):
        # Embed positions
        length = token_ids.shape[-1]
        positions = ops.arange(0, length, dtype="int32")
        positions_vectors = ops.take(self.position_embed, positions, axis=0)
        # Embed tokens
        token_ids = ops.cast(token_ids, dtype="int32")
        token_vectors = ops.take(self.token_embed, token_ids, axis=0)
        # Sum both
        embed = token_vectors + positions_vectors
        # Normalize embeddings
        power_sum = ops.sum(ops.square(embed), axis=1, keepdims=True)
        return embed / ops.sqrt(ops.maximum(power_sum, 1e-7))
```

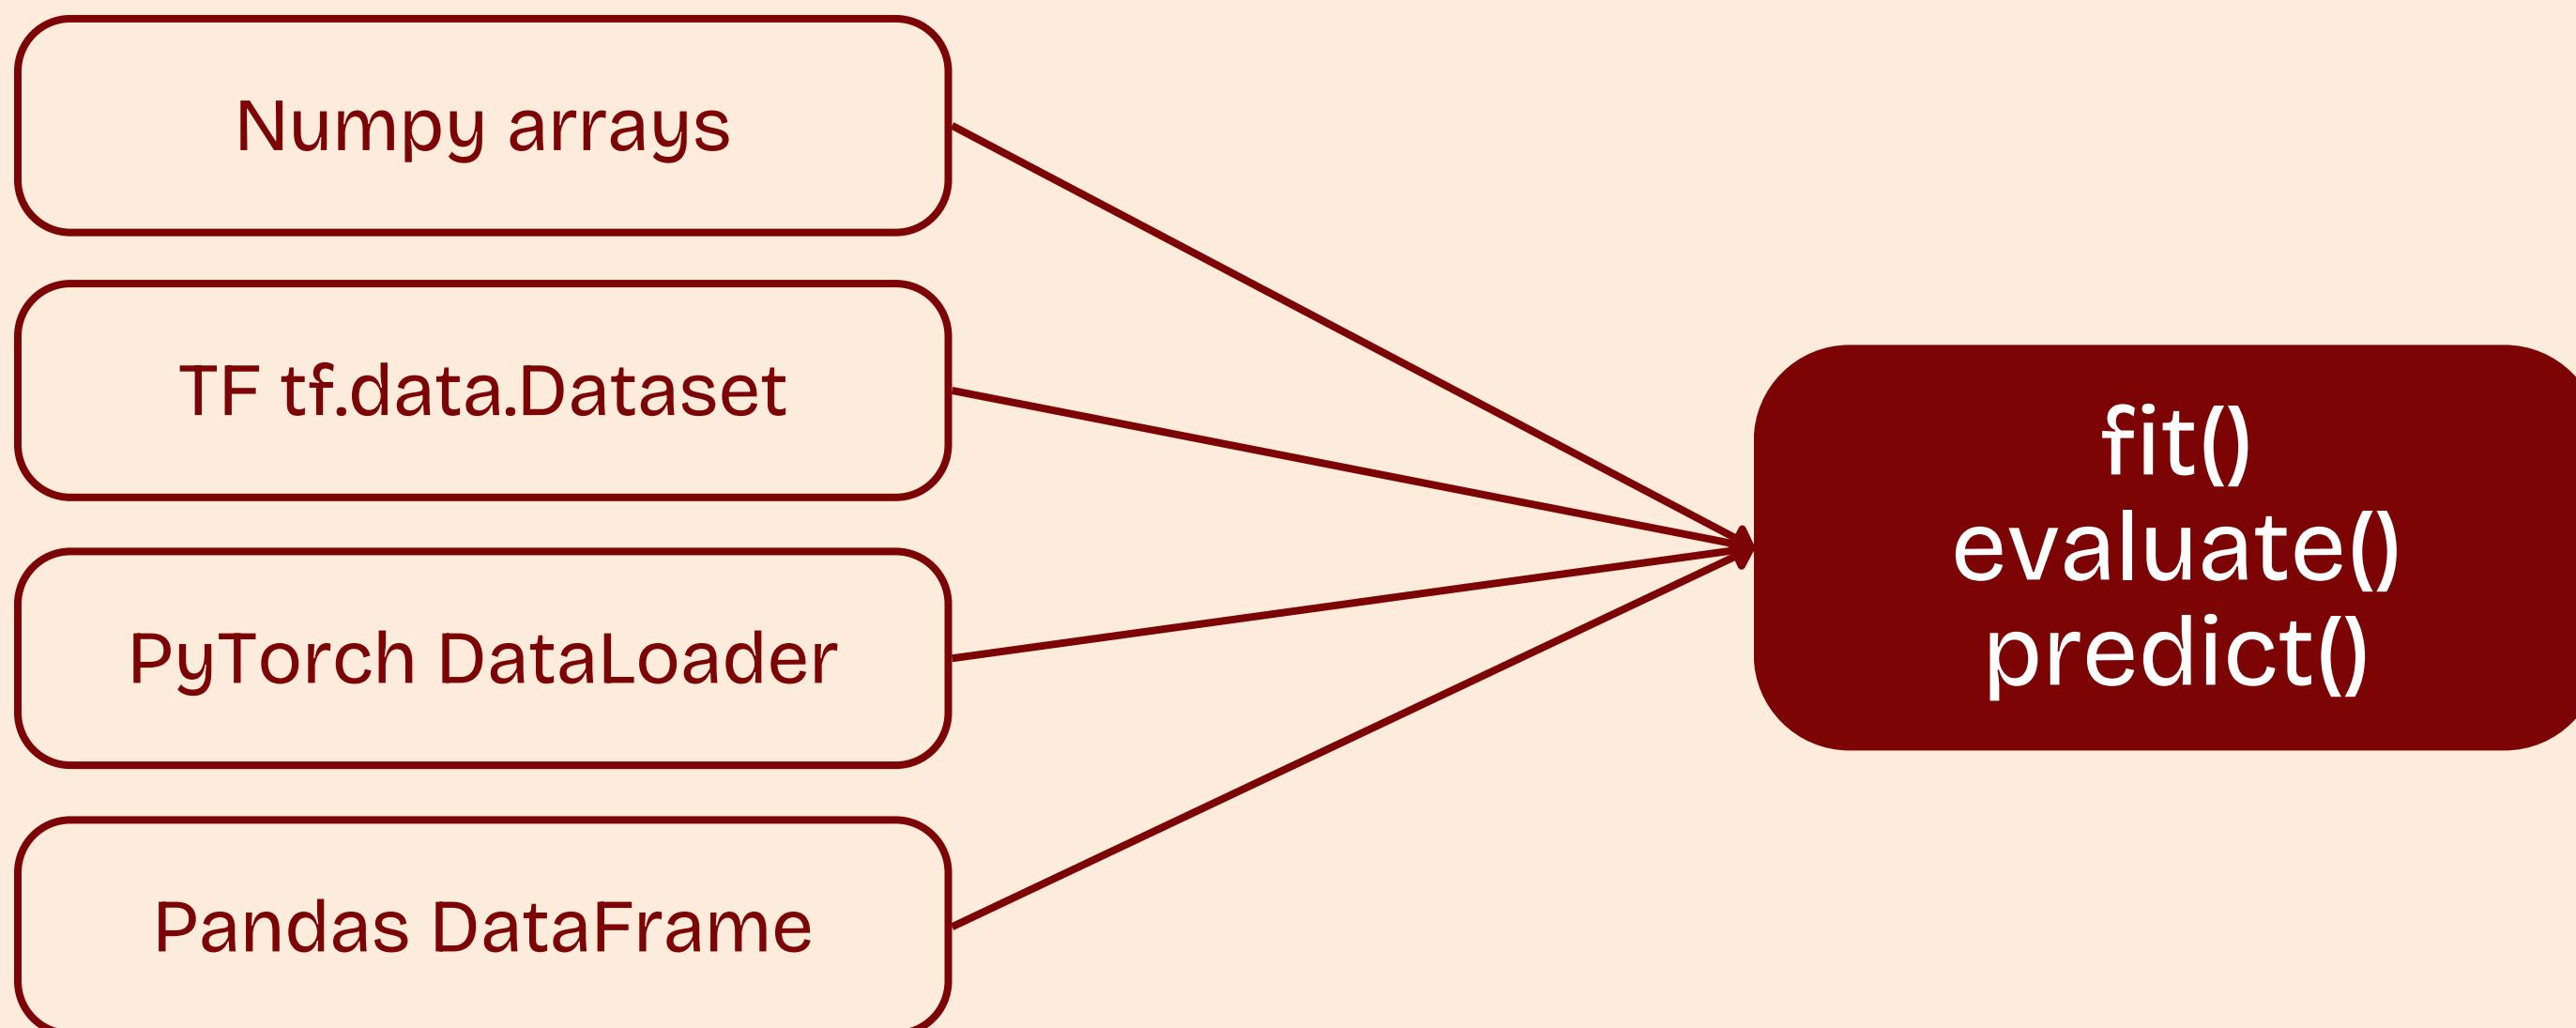
```
import os
os.environ["KERAS_BACKEND"] = "jax"
import keras
```

```
import os
os.environ["KERAS_BACKEND"] = "tensorflow"
import keras
```

```
import os
os.environ["KERAS_BACKEND"] = "torch"
import keras
```

You just have to specify which backend to run the code with!

Cross-framework data pipelines



This helps leverage the backends easily!

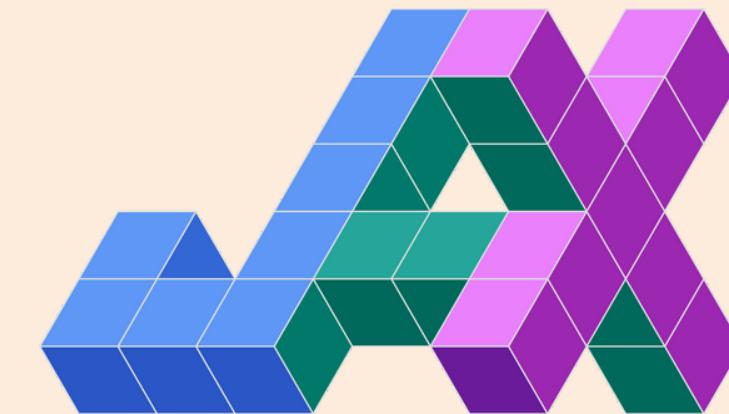


PyTorch
geometric

OpenMined

Skorch

Development



Large scale TPU training

Large scale parallelism

Training



TensorFlow

TF-Serving

TF.js

TFLite

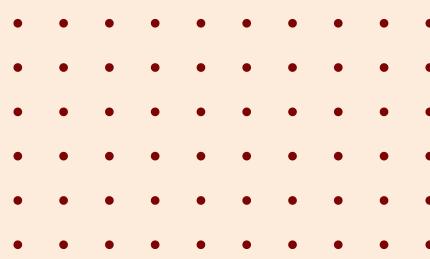
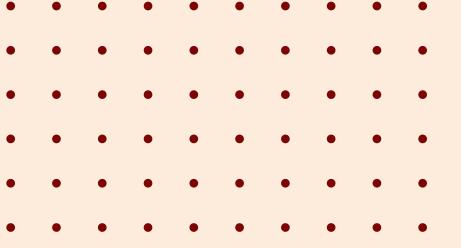
Deployment

03

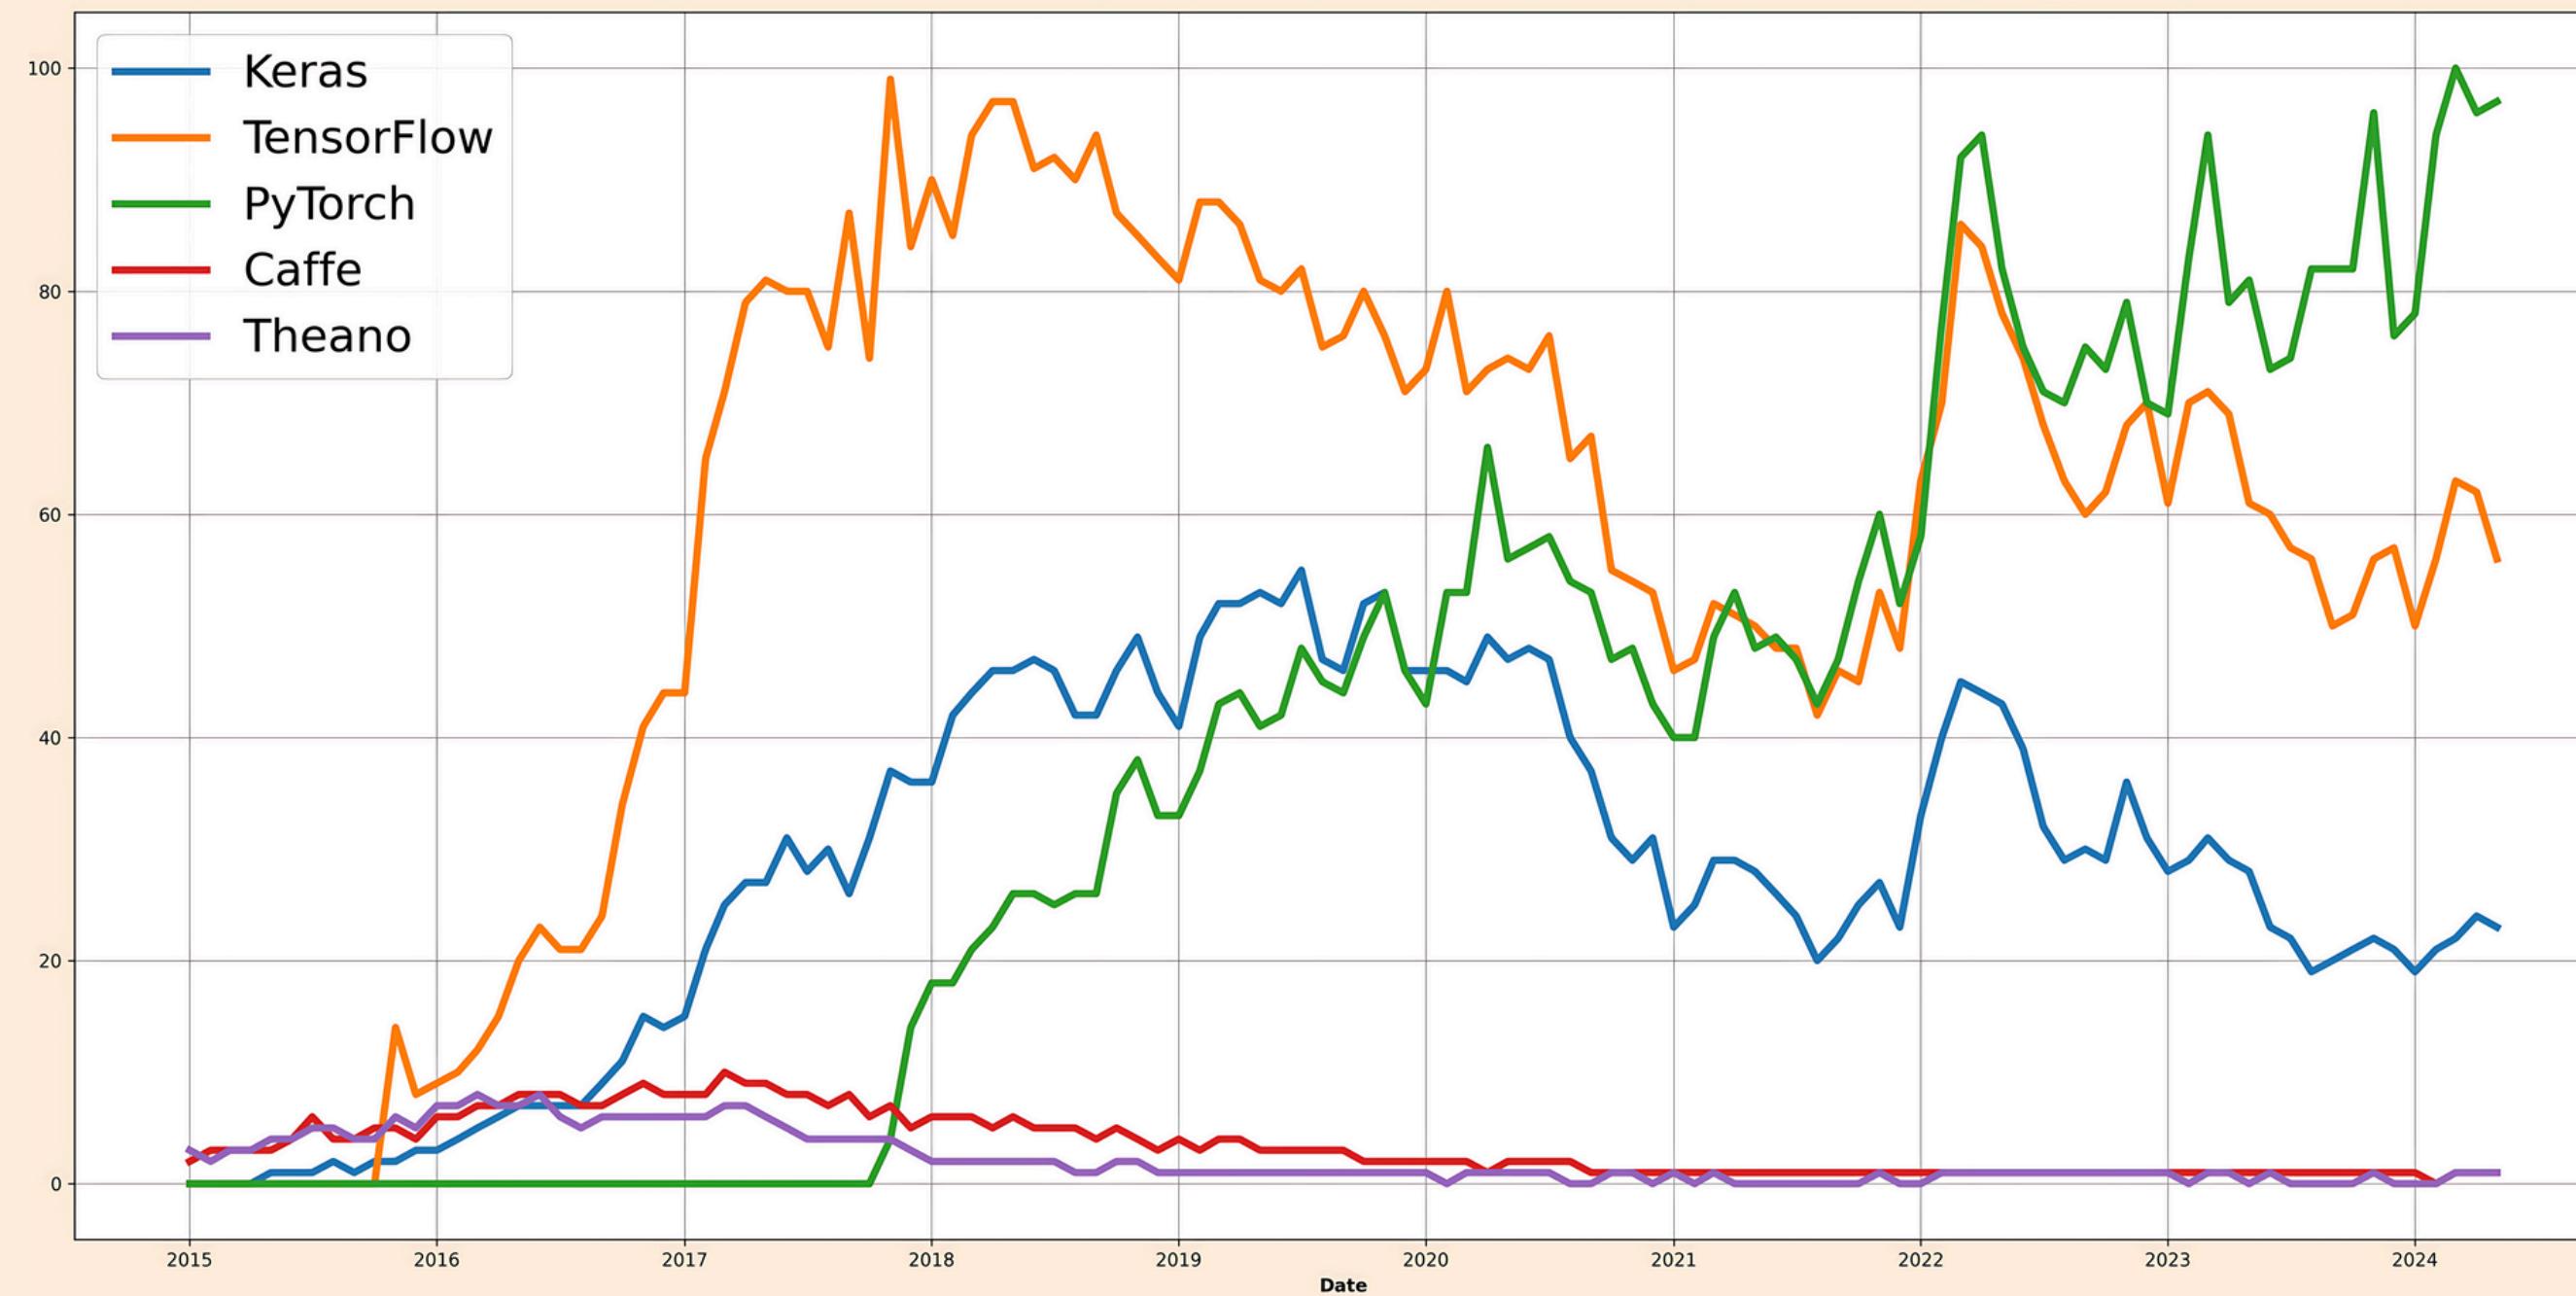
Demonstration

04

Conclusion



Level of interest over time of some deep learning libraries





Keras

VS

PyTorch



TensorFlow

Developed by Google Brain
Released in 2015
Supply Tensors library
High & Low-level DL framework

Developed by Facebook
Released in 2016
Based on Torch library
Low-level DL framework

Comparison

	Keras	TensorFlow	Pytorch
Level	High	High, Low	Low
Written in	Python	C++, CUDA, Python	Python, C++
Ease of use	Very high	Low	High
Development speed	Faster		Slower
Flexibility	Lower		Higher

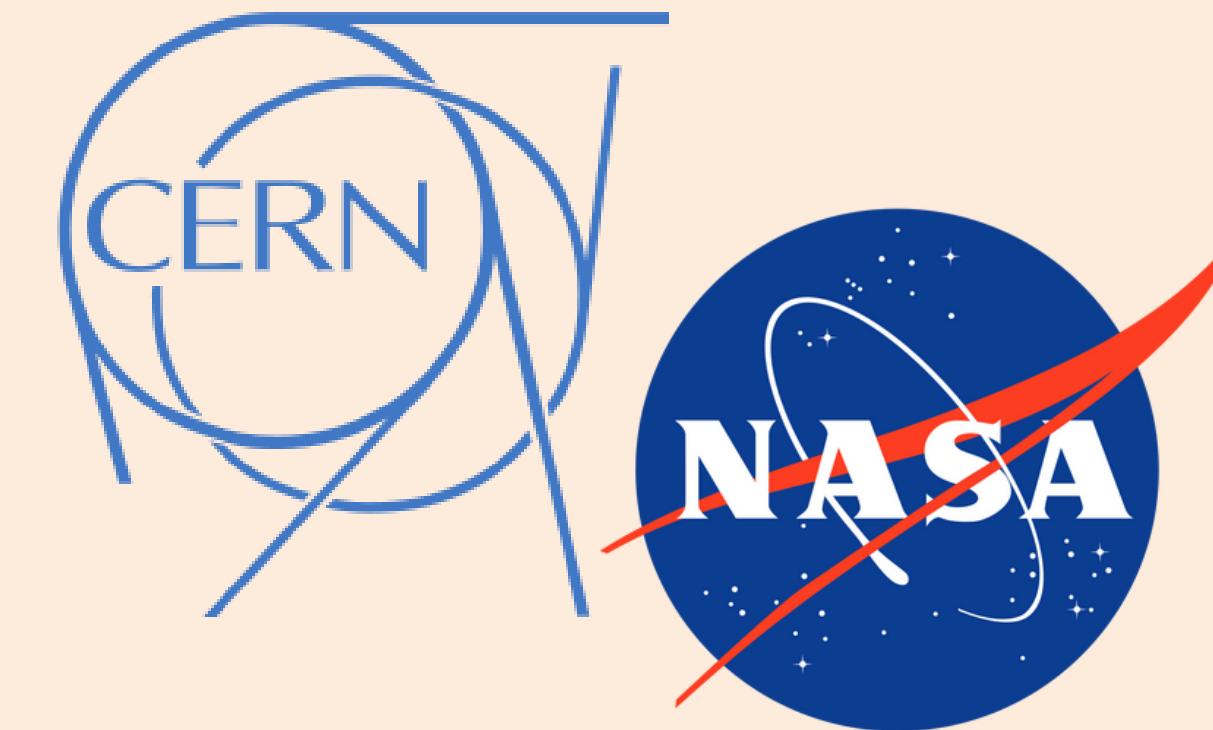
When to use

	Keras	TensorFlow	Pytorch
Beginners	✓		✓
Rapid prototyping	✓		✓
Flexibility required		✓	✓
Scalability required		✓	
Deployment required		✓	

Keras' Application

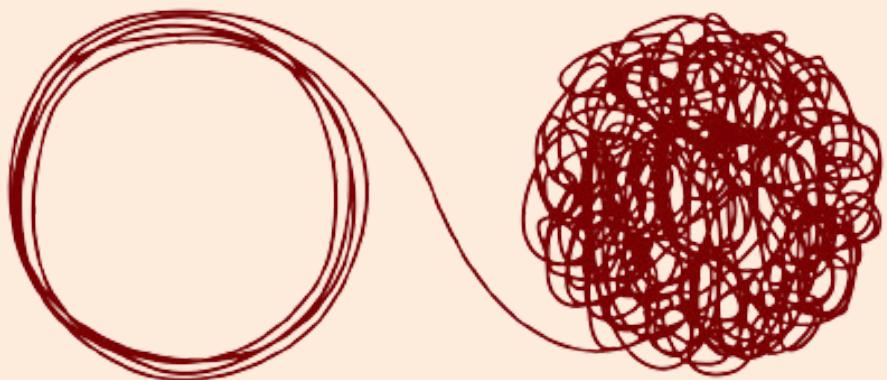
 YouTube
Uber

Industrial Applications



Scientific Organizations

**Thank you
for your attention!**



Q&A